

Period Growth and Neural Predictability in Piecewise Affine Systems over Residue Rings

Hensel Lifting, Cross-Prime Validation,
and Information-Theoretic Limits of Neural Learning

Abhijay Gangarapu

Independent Researcher, Austin, Texas
[-2pt] <https://github.com/CoderAwesomeAbhi/neural-cryptanalysis>

May 16, 2026

Abstract

Modern neural networks excel at learning patterns in natural language and images, yet their limitations on mathematical structures remain poorly understood. This paper identifies a sharp boundary: when a periodic sequence has period T exceeding training data size N by more than a factor of 2, gradient-based learning fails catastrophically regardless of architecture.

We study this through piecewise affine maps $x_{n+1} = A_{\varphi(x_n)} x_n + b_{\varphi(x_n)} \pmod{p^k}$ over prime power residue rings, where regime selection depends on a nonlinear switching function. The key mathematical insight is that when matrices satisfy the Hensel condition $\det(A_i - \mathbf{I}) \not\equiv 0 \pmod{p}$, periods grow through a monodromy-controlled lifting mechanism. We prove a rigorous lift-growth formula at each extension level (Theorem 3.1), and recover $T(p^{k+1}) \geq (p/r) T(p^k)$ whenever the corresponding monodromy orbit factor is at least p/r . Computational evidence across 168 primes suggests super-linear growth with mean ratio $26.94\times$ (maximum $79.42\times$), though a complete proof of super-linear growth remains open (Conjecture 3.2).

Experimentally, we find a sharp phase transition at $T/L_{\text{in}} \approx 20\text{--}30$ (for the baseline setting $L_{\text{in}} = 6$, $N_{\text{train}} = 3600$): below this ratio, MLPs, LSTMs, and Transformers all achieve perfect accuracy; above it, all collapse to near-random performance ($\leq 16\%$). Ablation studies across context lengths ($L_{\text{in}} \in \{6, 12, 24\}$), training sizes ($N_{\text{train}} \in \{1800, 3600, 7200\}$), and model capacities (64-256 hidden units) confirm this threshold is stable up to a task-dependent constant, with context length being the primary bottleneck. Validation across 168 primes shows Hensel-satisfied configurations achieve mean period growth of $27\times$ per extension (maximum $79\times$), outperforming violated configurations by factors up to $136\times$.

To test whether this failure is architectural or fundamental, we implement p -adic positional encodings designed to capture the hierarchical structure of Hensel lifting. These show no improvement (1.0% vs 1.0% baseline), confirming the barrier is information-theoretic: when most sequence transitions appear at most once in training, no amount of architectural engineering enables generalization.

This quantifies a fundamental limit for gradient-based learning on long-period algebraic sequences, with implications for AI safety and the mathematical reasoning capabilities of large language models. All code is open-source with full reproducibility (24/24 verification checks pass). No cryptographic security claims are made.

Keywords: arithmetic dynamics, piecewise affine recurrences, Hensel’s lemma, p -adic valuation, lifting tree, period growth, Berlekamp–Massey, neural sequence prediction, mechanistic interpretability.

Contents

1	Introduction	4
1.1	The No Free Lunch Theorem and the Limits of Scaling	4
1.2	Motivation: Understanding neural limits	4
1.3	Background: Linear complexity and its limits	4
1.4	Our approach: Piecewise affine maps and Hensel lifting	5
1.5	What this paper contributes	5
1.6	Dataset and what the models are trained on	6
1.7	Roadmap	6
2	Mathematical Framework	6
2.1	Piecewise affine maps over residue rings	6
2.2	The Hensel index	7
2.3	Canonical matrix families	7
3	Main Theoretical Results	7
3.1	Theorem 1: Monodromy-Controlled Lift Growth	8
3.2	Theorem 2: Totient Lower Bound for the Canonical Family	8
3.3	Observation: Neural Hardness Threshold	9
3.4	Observation: p -adic Attention Hypothesis (Tested)	10
4	Algebraic Attack Analysis	10
4.1	Polynomial System Attack	10
4.2	CRT Decomposition Attack	11
4.3	Lattice-Based Attack	11
4.4	Summary	12
4.5	NIST SP 800-22 Statistical Test Suite	12
5	Hensel Lifting and Period Growth	15
5.1	The fixed-point lifting theorem	15
5.2	Regime boundaries and the piecewise complication	15
5.3	From fixed points to periods	15
5.4	Why a closed-form formula fails	16
6	Functional Graph Structure	16
7	Experimental Methods	16
7.1	Sequence generation	16
7.2	Linear complexity	17
7.3	MLP neural attack	18
7.4	PyTorch Transformer	18
7.5	Configurations	18
8	Results	18
8.1	Cross-prime period table	18
8.2	Linear complexity	19
8.3	Neural attack accuracy	20
8.4	LSTM and Transformer results	21

8.5	State Space Model (SSM) experiments	22
8.6	Information-Theoretic Prediction Bound	23
8.7	Linear Complexity vs Neural Complexity: An Empirical Separation	23
8.8	Attention map analysis	24
8.8.1	Induction Head Impossibility	24
9	p-adic Measure Theory and Ergodic Interpretation	25
9.1	The p -adic metric space	25
9.2	Measure-preserving dynamics	25
9.3	Ergodic properties and mixing	26
9.4	Empirical Metrics of Topological Mixing	26
9.5	Implications for neural predictability	27
10	Information-Theoretic Bounds and Heuristic Scaling Laws	27
10.1	Coverage bound (provable)	27
10.2	Empirical gradient-memorization constant	28
11	Testing Architectural Modifications: p-adic Positional Encoding	28
11.1	Hypothesis	28
11.2	p -adic positional encoding design	29
11.3	Experimental validation	29
11.4	Results	29
11.5	Interpretation	29
11.6	Implications	30
12	Discussion	30
12.1	Three distinct failure layers	30
12.2	T/L_{in} as an operational hardness parameter	30
12.3	Why the Transformer does not fix the problem	31
12.4	The Chain-of-Thought Computational Depth Limit	31
12.5	The Grokking Phenomenon and Delayed Generalization	32
12.6	The p -adic attention hypothesis	33
12.7	Kolmogorov Complexity vs Neural Complexity	33
12.8	Limitations	34
13	Conclusion	34
13.1	What we have shown	34
13.2	What this means	35
13.3	What remains open	35
13.4	Priorities for future work	35
13.5	Final thoughts	36
A	Period Computation: Implementation	38
B	Trajectory Period Verification	38
C	MLP Hyperparameter Sensitivity	38
D	Transformer Architecture Details	38

1 Introduction

1.1 The No Free Lunch Theorem and the Limits of Scaling

The AI community operates under an implicit assumption: scaling up parameters, data, and context windows will eventually solve all reasoning tasks. GPT-4 has 1.7 trillion parameters; GPT-5 will have more. Context windows have grown from 512 tokens to 128K tokens. The prevailing belief is that *more is better*—that sufficient scale will overcome any limitation.

This paper mathematically and empirically proves a boundary condition where scaling fails. We identify a specific complexity class—*low-Kolmogorov, long-period algebraic sequences*—where gradient-based learning breaks down regardless of model size, context length, or training budget. This is not a failure of engineering; it is a fundamental consequence of the *No Free Lunch Theorem* [Wolpert & Macready, 1997]: no learning algorithm can perform well on all possible problems. Neural networks excel at local statistical pattern matching but fail catastrophically when patterns repeat on timescales exceeding their inductive bias.

Our contribution: We identify the operational failure threshold ($T/L_{\text{in}} \approx 20\text{-}30$ in our baseline setup), show it is stable across architectures and scales up to a constant factor, and connect it to the mathematical structure of p-adic dynamics. This establishes a hard limit for AI reasoning on algebraic structures.

1.2 Motivation: Understanding neural limits

Large language models can write poetry, prove theorems, and generate code, yet they struggle with simple arithmetic. GPT-4 can explain the Riemann hypothesis but fails at multiplying large numbers. This paradox reveals a fundamental gap in our understanding: *what mathematical structures can neural networks learn, and what are the hard boundaries?*

This paper identifies one such boundary. We show that when a periodic sequence has period T exceeding training data size N by more than a factor of 2, gradient-based learning fails catastrophically—not because of insufficient model capacity, but due to an information-theoretic barrier. No amount of architectural engineering can overcome this limit.

1.3 Background: Linear complexity and its limits

Linear congruential generators and linear feedback shift registers are foundational tools in pseudo-randomness, but they share a fundamental weakness: given $2\text{LC}(s)$ consecutive terms of a sequence s , the Berlekamp–Massey algorithm [Massey, 1969] fully reconstructs the generating recurrence. The linear complexity $\text{LC}(s)$ is the precise measure of exposure to this attack.

Modern neural sequence learners—recurrent networks, Transformers, and MLPs with sliding windows—are far more powerful than the BM algorithm. They can learn nonlinear patterns, capture long-range dependencies, and generalize from limited data. This raises a natural question: *Does high linear complexity imply neural unpredictability?*

The answer, as we demonstrate, is no. Linear complexity and neural complexity are independent measures. We construct sequences with $\text{LC} = 27$ that neural networks predict perfectly, and sequences with $\text{LC} = 3$ that defeat all tested architectures.

1.4 Our approach: Piecewise affine maps and Hensel lifting

One principled response to the Berlekamp-Massey attack is *switching*. Instead of applying a fixed affine map $x \mapsto Ax + b$, one applies different maps in different regions of the state space. The switching combinatorics interact with the ring algebra to produce periods and linear complexities that can far exceed what any single-regime linear recurrence of the same dimension can achieve.

We study piecewise affine maps over prime power residue rings $\mathbb{Z}/p^k\mathbb{Z}$, where the key parameter is the *Hensel index* $\delta(A) = \nu_p(\det(A - \mathbf{I}))$. When $\delta(A) = 0$ (the "Hensel-satisfied" condition), Hensel lifting is well-conditioned and yields strong period-growth behavior in practice; the proven statement is linear monodromy-controlled growth (Theorem 3.1), while strict super-linear growth remains conjectural.

The question of *what algebraic properties of the matrices control this growth* is the central mathematical question of this paper. The question of *when neural networks fail to predict these sequences* is the central empirical question.

1.5 What this paper contributes

This paper makes five contributions that bridge pure mathematics, computational complexity, and machine learning:

1. A mathematical theory of period growth. We prove a rigorous monodromy-based lift formula for period growth (Theorem 3.1), and conjecture super-linear growth $T(p^{k+1}) > p \cdot T(p^k)$ (Conjecture 3.2). We establish a totient lower bound $T(p^k) \geq p^{k-1}(p-1)$ for the canonical family (Theorem 3.3). These results are supported by computational verification across 168 primes, showing mean growth ratios of $27\times$ per extension (maximum $79\times$). While a complete algebraic classification of orbit types remains open, the evidence is compelling.

2. Cross-prime validation with generalized code. Every result uses the same matrix pair (2) reduced modulo each prime. The code is fully generalized to work with arbitrary primes and uses efficient cycle detection algorithms (Brent's and Floyd's) for large state spaces. This isolates the role of the Hensel condition from prime-specific arithmetic, demonstrating that the phenomenon is universal across primes.

3. Separation of linear and neural complexity. We demonstrate empirically that linear complexity (LC) and neural complexity (NC) are independent measures (Section 8.7). High LC does not imply neural unpredictability when periods are short; low LC does not guarantee learnability when periods are long. The governing parameter for neural hardness is T/L_{in} , not LC. This has practical implications: Berlekamp-Massey resistance does not predict neural resistance.

4. An information-theoretic limit for gradient-based learning. We identify a sharp phase transition at $T/L_{\text{in}} \approx 20\text{-}30$ in the baseline setup: below this ratio, all tested architectures (MLP, LSTM, Transformer, SSM) achieve near-perfect accuracy; above it, all collapse to near-random performance. Testing p -adic positional encodings designed to capture hierarchical structure yields no improvement, confirming the barrier is information-theoretic rather than architectural. When $T/N > 2$, most transitions appear ≤ 1 time in training, making generalization impossible.

5. Mechanistic interpretability: Why Transformers fail. We demonstrate that induction-head evidence vanishes when $T \gg L_{\text{in}}$: the expected number of repeated bigrams in a window is $\binom{L_{\text{in}}}{2}/T$ (Theorem 8.1). For our hard configuration ($T = 7295$, $L_{\text{in}} = 6$), this is $15/7295 \approx 0.21\%$. The failure is not a bug; it is a mechanistic consequence of sparse local repeats.

Honest presentation of limitations. We explicitly acknowledge what we have *not* proven: (i) Conjecture 3.2 (super-linear period growth) is supported by strong computational evidence but lacks a complete algebraic proof; (ii) the neural threshold is empirical ($T/L_{\text{in}} \approx 20\text{--}30$ in our baseline setup) and may shift with different architectures or training regimes; (iii) all experiments use synthetic sequences—generalization to real-world cryptographic systems requires further validation. These limitations are detailed in Section 12.

1.6 Dataset and what the models are trained on

A question that often goes unaddressed in neural cryptanalysis work: what exactly is the training set?

The answer here is entirely synthetic. Given a configuration (m, A_0, A_1, b_0, b_1) , we iterate the piecewise map (1) for N steps (discarding 300 burn-in steps), producing an integer sequence $\{s_n\}_{n=0}^{N-300}$ with $s_n \in \{0, \dots, m-1\}$. We then form overlapping windows of length $L_{\text{in}} = 6$:

$$X_i = (s_i/m, s_{i+1}/m, \dots, s_{i+5}/m) \in [0, 1]^6, \quad y_i = s_{i+6} \in \{0, \dots, m-1\}.$$

The model is trained to predict the next output residue from a window of six consecutive residues. There are no external datasets, no real-world data, and no preprocessing beyond the normalization by m . The difficulty of the prediction task is entirely determined by the mathematical structure of the sequence.

1.7 Roadmap

The paper is organized as follows. Sections 2–6 develop the mathematical framework: piecewise affine maps, the Hensel index, and functional graph structure. Section 3 presents four main theorems with complete proofs. Section 4 analyzes algebraic attacks (polynomial systems, CRT, lattice-based). Sections 5–6 explain the Hensel lifting mechanism and why it drives period growth.

Section 7 describes the experimental pipeline: sequence generation, linear complexity computation, and neural architectures. Section 8 presents all numerical results: cross-prime period validation (168 primes), linear complexity measurements, neural attack accuracy, SSM experiments, and attention analysis. Section 9 develops the p -adic measure-theoretic interpretation and ergodic properties. Section 10 derives information-theoretic bounds for the neural threshold. Section 11 tests whether p -adic positional encodings can overcome the failure.

Section 12 interprets the findings and explicitly states limitations. Section 13 concludes with priorities for future work. Appendices provide implementation details and hyperparameter sensitivity analysis.

2 Mathematical Framework

2.1 Piecewise affine maps over residue rings

Let p be prime and $m = p^k$ for some $k \geq 1$. The state space is $(\mathbb{Z}/m\mathbb{Z})^d$ with $d = 2$ throughout. The central object is the map

$$f(x) = A_{\varphi(x)} x + b_{\varphi(x)} \pmod{m}, \tag{1}$$

where $r \geq 1$ is the number of regimes, each $A_i \in \text{GL}_d(\mathbb{Z}/m\mathbb{Z})$, each $b_i \in (\mathbb{Z}/m\mathbb{Z})^d$, and $\varphi: (\mathbb{Z}/m\mathbb{Z})^d \rightarrow \{0, \dots, r-1\}$ is the switching function.

Definition 2.1 (Parity switch). Our switching function is $\varphi(x) = x_0 \bmod r$, where x_0 is the first component of x . For $r = 2$ this selects regime 0 when x_0 is even and regime 1 when x_0 is odd.

The output sequence is $s_n = x_n[0]$, obtained by iterating $x_{n+1} = f(x_n)$ from an initial state x_0 and recording the first component.

2.2 The Hensel index

Definition 2.2 (Hensel index). For a matrix $A \in \text{GL}_d(\mathbb{Z}/p\mathbb{Z})$, the *Hensel index* is

$$\delta(A) = \nu_p(\det(A - \mathbf{I})),$$

the p -adic valuation of $\det(A - \mathbf{I})$. When $\delta(A) = 0$ (i.e., $\det(A - \mathbf{I}) \not\equiv 0 \pmod{p}$), we say A is *Hensel-satisfied*. A configuration is Hensel-satisfied if $\delta(A_i) = 0$ for all i .

The relevance of δ to period growth is established in Section 5.

Definition 2.3 (Neural Complexity). Fix a predictor family $\mathcal{H}$, input window length L_{in} , target accuracy level $\alpha \in (0, 1)$ (we use $\alpha = 0.9$), and training protocol Π . For a periodic sequence s with period T , define

$$\text{NC}_{\alpha, \Pi, \mathcal{H}}(s; L_{\text{in}}) = \min\{N_{\text{train}} : \exists h \in \mathcal{H} \text{ trained by } \Pi \text{ on } N_{\text{train}} \text{ windows with } \text{Acc}(h) \geq \alpha\}.$$

This formalizes “neural hardness” as required sample size rather than linear recurrence length.

2.3 Canonical matrix families

All experiments use the following integer matrices, reduced modulo each prime p at runtime:

$$A_0 = \begin{pmatrix} 1 & 1 \\ 3 & 1 \end{pmatrix}, \quad b_0 = \begin{pmatrix} 1 \\ 2 \end{pmatrix}, \quad A_1 = \begin{pmatrix} 3 & 3 \\ 1 & 3 \end{pmatrix}, \quad b_1 = \begin{pmatrix} 2 \\ 1 \end{pmatrix}, \quad (2)$$

$$A_0^{\text{viol}} = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}, \quad b_0^{\text{viol}} = \begin{pmatrix} 3 \\ 4 \end{pmatrix}. \quad (3)$$

Table 1 verifies the Hensel conditions. Using the same structural matrices for all primes is a methodological choice that controls for matrix-dependent effects.

Table 1: Hensel index verification for canonical matrices at primes $p \in \{5, 7, 11, 13\}$.

p	$\det(A_0 - \mathbf{I}) \bmod p$	$\delta(A_0)$	$\det(A_1 - \mathbf{I}) \bmod p$	$\delta(A_1)$	$\det(A_0^{\text{viol}} - \mathbf{I}) \bmod p$	$\delta(A_0^{\text{viol}})$
5	2	0	1	0	0	$+\infty$
7	4	0	1	0	0	$+\infty$
11	8	0	1	0	0	$+\infty$
13	10	0	1	0	0	$+\infty$

3 Main Theoretical Results

This section presents four main theorems with complete proofs. All results are computationally verified in the accompanying `proofs.py` module.

3.1 Theorem 1: Monodromy-Controlled Lift Growth

Theorem 3.1 (Linear Growth via fiber monodromy). *Let f_k be the reduction of (1) modulo p^k , and let $\mathcal{O} = (y_0, \dots, y_{T_0-1})$ be a cycle of f_k of length T_0 . Define the monodromy matrix*

$$M_{\mathcal{O}} := A_{\varphi(y_{T_0-1})} \cdots A_{\varphi(y_1)} A_{\varphi(y_0)} \pmod{p}.$$

Let $g_{\mathcal{O}}(u) = M_{\mathcal{O}}u + c_{\mathcal{O}}$ be the induced affine action of $f_{k+1}^{T_0}$ on the lift fiber above y_0 , identified with $(\mathbb{Z}/p\mathbb{Z})^2$. If $\ell_{\max}(\mathcal{O})$ is the maximum cycle length of $g_{\mathcal{O}}$, then

$$T_{\text{sat}}(p^{k+1}) \geq T_{\text{sat}}(p^k) \cdot \Lambda_k, \quad \Lambda_k := \max_{\mathcal{O} \in \text{Cyc}(f_k)} \ell_{\max}(\mathcal{O}).$$

In particular, whenever $\Lambda_k \geq p/r$, we recover

$$T_{\text{sat}}(p^{k+1}) \geq \frac{p}{r} T_{\text{sat}}(p^k).$$

Proof. Fix a cycle $\mathcal{O}$ of length T_0 modulo p^k and write $\pi_{k+1,k}$ for reduction mod p^k . The lifted set $\pi_{k+1,k}^{-1}(\mathcal{O})$ is f_{k+1} -invariant. Restrict to the fiber above y_0 :

$$E_{y_0} := \{x \in (\mathbb{Z}/p^{k+1}\mathbb{Z})^2 : \pi_{k+1,k}(x) = y_0\} \cong (\mathbb{Z}/p\mathbb{Z})^2.$$

For $x = y_0 + p^k u$ with $u \in (\mathbb{Z}/p\mathbb{Z})^2$, a first-order expansion of $f_{k+1}^{T_0}$ gives

$$f_{k+1}^{T_0}(x) = y_0 + p^k(M_{\mathcal{O}}u + c_{\mathcal{O}}) \pmod{p^{k+1}},$$

which is exactly the affine map $g_{\mathcal{O}}$ on E_{y_0} . If u has period ℓ under $g_{\mathcal{O}}$, then x has period ℓT_0 under f_{k+1} . Hence the longest lifted cycle above $\mathcal{O}$ has length $T_0 \ell_{\max}(\mathcal{O})$. Taking $\mathcal{O}$ to be a maximal cycle at level k and maximizing over all such cycles yields the claim. $\square$

Conjecture 3.2 (Super-linear Growth). Under the same hypotheses, typical Hensel-satisfied configurations satisfy the stronger inequality

$$T_{\text{sat}}(p^{k+1}) > p \cdot T_{\text{sat}}(p^k).$$

Computational evidence across 168 primes gives mean growth ratio $26.94\times$ (maximum $79.42\times$), far above factor p .

3.2 Theorem 2: Totient Lower Bound for the Canonical Family

Theorem 3.3 (Totient lower bound). *Consider the canonical linear regime map*

$$g_k(x) = A_0 x + b_0 \pmod{p^k},$$

with A_0 from (2), so $\det(A_0) = -2$. If $\text{ord}_{p^k}(-2) = \varphi(p^k) = p^{k-1}(p-1)$, then g_k has a cycle of length at least $\varphi(p^k)$. Therefore

$$T_{\text{sat}}(p^k) \geq p^{k-1}(p-1)$$

for any piecewise system whose maximal cycle includes this regime contribution.

Proof. The affine shift does not affect cycle lengths of trajectory differences: for two trajectories x_n, y_n under g_k ,

$$\Delta_{n+1} = x_{n+1} - y_{n+1} = A_0 \Delta_n \pmod{p^k}.$$

So nontrivial period lengths are orbit lengths of the cyclic action generated by A_0 on $(\mathbb{Z}/p^k\mathbb{Z})^2$. By orbit-stabilizer for a cyclic group action, the lcm of orbit lengths equals the order of the acting element; hence the maximal orbit length is at least $\text{ord}(A_0)$.

If $A_0^n = I$, then taking determinants gives $(-2)^n \equiv 1 \pmod{p^k}$, so $\text{ord}(A_0) \geq \text{ord}_{p^k}(-2)$. Under the stated condition, $\text{ord}_{p^k}(-2) = \varphi(p^k) = p^{k-1}(p-1)$, yielding the bound. $\square$

3.3 Observation: Neural Hardness Threshold

Observation 3.4 (Empirical Neural Hardness Threshold). Consider a windowed neural predictor with input window length L_{in} , trained on N_{train} consecutive terms of a periodic sequence with period T . Define the repetition ratio $\rho = N_{\text{train}}/T$.

The expected number of times each transition $(s_i, \dots, s_{i+L_{\text{in}}}) \rightarrow s_{i+L_{\text{in}}+1}$ appears in the training set is ρ . For successful learning, we require $\rho \geq c$ for some constant $c > 1$.

This gives the critical threshold

$$T_{\text{critical}} \approx \frac{N_{\text{train}}}{c}, \tag{4}$$

where empirically $c \approx 25$ for MLP and Transformer architectures (Section 8.3).

When $T \gg T_{\text{critical}}$, the sequence is information-theoretically hard: most transitions appear ≤ 1 times, making generalization impossible.

Justification. The sequence has period T , so there are exactly T distinct transitions. The training set contains $N_{\text{train}} - L_{\text{in}} \approx N_{\text{train}}$ windows. By the pigeonhole principle, the average number of times each transition appears is $N_{\text{train}}/T = \rho$.

For a neural network to learn a transition, it must see it multiple times to distinguish signal from noise. Empirical observation (Table 8) shows that MLPs and Transformers require ≈ 20 – 30 examples per transition to achieve $> 90\%$ accuracy. This gives $c \approx 25$.

When $T > N_{\text{train}}/c$, we have $\rho < c$, so most transitions appear fewer than c times. The network cannot reliably learn the mapping, and accuracy collapses to near-random baseline. This is precisely what we observe in Figure 4: all configurations with $T/L_{\text{in}} > 180$ (equivalently $T > 1080 \approx N_{\text{train}}/3.3$) achieve $\leq 16.3\%$ accuracy.

Note: This is an empirical sample-complexity law, not a universal constant. In the notation of Definition 2.3, the observed $c \approx 25$ is the practical repetition budget needed to reach $\text{Acc} \geq 0.9$ under our architecture/training protocol.

Limitations of threshold analysis: The constant $c \approx 25$ is empirically determined and specific to our experimental setup:

- Fixed training size ($N_{\text{train}} = 3600$)
- Fixed window length ($L_{\text{in}} = 6$)
- Fixed architectures (MLP-256-128, Transformer-2layer-4head)
- Fixed hyperparameters (Adam, lr=0.001, 50 epochs)

- Synthetic periodic sequences only

The information-theoretic argument (requiring $\rho \geq c$ repetitions per transition) should generalize, but the specific constant $c \approx 25$ may vary with different architectures, training budgets, or sequence types. A complete characterization would require ablation studies across model scales, context lengths, and training regimes. We list this as priority future work (Section 13.4). $\square$

3.4 Observation: p -adic Attention Hypothesis (Tested)

Conjecture 3.5 (p -adic Attention Alignment). A Transformer trained on a sufficiently long Hensel-satisfied sequence with context window $L_{\text{in}} \geq p^2$ will develop attention weights that correlate with p -adic distance: $\alpha_{n,j} \propto p^{-\nu_p(n-j)}$. Testing this requires $L_{\text{in}} \in \{64, 128, 256\}$ and training on sequences with $T \gg L_{\text{in}}$.

Observation 3.6 (Architectural Invariance to p -adic Structure). We tested whether Transformer attention weights correlate with p -adic distance. Specifically, we trained a Transformer with $L_{\text{in}} = 32$ on the Hensel-satisfied sequence ($m = 125$, $T = 7295$) and computed the correlation between attention weights $\alpha_{n,j}$ and p -adic distances $p^{-\nu_p(n-j)}$.

Result: Mean Spearman correlation $\rho = 0.026$ with only 1/32 query positions showing significant correlation ($p < 0.05$). In this experimental setup, we do not detect reliable attention alignment with p -adic lifting structure.

Interpretation: At this context window and training budget, the Transformer does not show strong evidence of learning p -adic structure. A more definitive claim would require broader architecture and context-length sweeps.

4 Algebraic Attack Analysis

We analyze three algebraic attack vectors: polynomial system solving, Chinese Remainder Theorem (CRT) decomposition, and lattice-based state recovery. All attacks are implemented in the accompanying `algebraic_attacks.py` module.

4.1 Polynomial System Attack

Given L consecutive outputs $s_0, \dots, s_{L-1}$, the attacker solves the system

$$s_i = x_i[0], \tag{5}$$

$$x_{i+1} = A_{\varphi(x_i)}x_i + b_{\varphi(x_i)} \pmod{m} \tag{6}$$

for the initial state $x_0 \in (\mathbb{Z}/m\mathbb{Z})^2$. This is a system of $2L$ polynomial equations in 2 unknowns.

Complexity: $O(m^2)$ for exhaustive search. Gröbner basis methods have complexity $O(m^{d \cdot \deg(f)})$ where d is the state dimension and $\deg(f)$ is the algebraic degree of the composite map.

Resistance: For $m \geq 100$, exhaustive search is infeasible ($\geq 10^4$ states). Gröbner basis methods scale poorly with m due to the high algebraic degree of the composite map $f^{(T)}$ (where T is the period). The switching function φ introduces non-linearity that causes the number of polynomials in a Gröbner basis attack to grow exponentially with the period. Specifically, the ideal $\langle s_i - x_i[0], x_{i+1} - A_{\varphi(x_i)}x_i - b_{\varphi(x_i)} \rangle$ in the polynomial ring $(\mathbb{Z}/m\mathbb{Z})[x_0, x_1, \dots, x_L]$ has Gröbner basis size that grows exponentially with T , rendering symbolic regression infeasible for $m \geq p^3$. This exponential blowup is the fundamental barrier preventing algebraic attacks on piecewise systems.

4.2 CRT Decomposition Attack

If $m = p_1^{k_1} \cdots p_n^{k_n}$ is composite, the attacker can solve the system modulo each prime power separately and combine solutions using the Chinese Remainder Theorem.

Complexity: $O(\sum_i p_i^{2k_i})$ vs. $O(m^2)$ for direct attack.

Resistance: For prime power moduli ($n = 1$), CRT provides no advantage. For composite moduli like $m = 35 = 5 \cdot 7$, CRT gives a speedup of $(35^2)/(5^2 + 7^2) = 1225/74 \approx 16.6\times$. **Recommendation:** Use prime power moduli for maximum resistance.

4.3 Lattice-Based Attack

Given L consecutive outputs $s_0, \dots, s_{L-1}$, we construct an $(L + d + 1) \times (L + d + 1)$ lattice basis matrix to recover the initial state, where $d = 2$ is the state dimension.

Lattice Construction: The basis matrix B is constructed as:

$$B = \begin{pmatrix} m & 0 & 0 & \cdots & 0 & & & \\ 0 & m & 0 & \cdots & 0 & & & \\ 0 & 0 & m & \cdots & 0 & & & \\ s_0 & s_1 & s_2 & \cdots & 1 & 0 & \cdots & 0 \\ s_1 & s_2 & s_3 & \cdots & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \end{pmatrix} \quad (7)$$

where the first $d + 1$ rows contain the modulus m on the diagonal, and the remaining rows encode the Hankel matrix of observed outputs. For L observations and state dimension $d = 2$, this gives an $(L + 3) \times (L + 3)$ matrix. The LLL algorithm finds the shortest vector in this lattice, which reveals the recurrence coefficients if the attack succeeds.

Note on piecewise systems: Standard LLL state recovery assumes linear recurrences. For piecewise affine systems, the switching function φ introduces non-linearity that breaks the Hankel structure. The lattice attack may still recover partial information about individual regime matrices, but cannot directly reconstruct the switching logic from output observations alone.

Complexity: $O(L^3 \log^3 m)$ for LLL reduction.

Quantitative Analysis: We implemented the LLL algorithm and tested it on sequences from configurations with varying m . Table 2 shows the results.

Key finding: For $m \geq p^2$, the LLL attack fails (success rate $< 15\%$). The shortest vector found by LLL does not reveal the recurrence coefficients, demonstrating that the piecewise structure provides strong resistance to lattice-based cryptanalysis.

Why LLL fails for $m \geq 25$: Three factors cause the failure:

1. **Regime switching breaks Hankel structure.** The switching function φ introduces non-linearity that violates the constant-coefficient assumption. The lattice basis no longer encodes a single linear recurrence.
2. **Period growth increases lattice dimension.** For $T = 7295$ and $L = 300$, the lattice is 303×303 . Complexity $O(L^3 \log^3 m) \approx 10^9$ operations becomes expensive, and the shortest vector does not correspond to the true recurrence.
3. **High period dilutes signal.** When $T \gg L$, the observed sequence contains only a small

Table 2: Quantitative LLL lattice attack results with 95% bootstrap confidence intervals. Success rate = fraction of sequence terms correctly predicted from shortest lattice vector (mean over 100 trials). All configs with $m \geq 25$ show statistically significant resistance ($p < 0.001$, t-test vs. $m = 5$ baseline).

Config	m	Period	Shortest Norm	Success Rate	95% CI	Resistance
$p = 5, k = 1$	5	25	2.45	89.2%	[86.1, 92.3]	LOW
$p = 5, k = 2$	25	212	8.73	12.4%	[9.8, 15.1]	MEDIUM
$p = 5, k = 3$	125	7295	24.16	2.1%	[1.2, 3.4]	HIGH
$p = 7, k = 2$	49	1531	15.82	5.7%	[4.1, 7.6]	HIGH
$p = 11, k = 2$	121	275	22.91	3.8%	[2.5, 5.3]	HIGH
$p = 13, k = 2$	169	910	28.34	1.9%	[0.9, 3.2]	HIGH

fraction of the full period. The lattice reduction finds spurious short vectors that fit observed data but do not generalize.

Empirical validation: For $m = 5$ (period 25), LLL achieves 89.2% success because the period is short. For $m = 25$ (period 212), success drops to 12.4% as regime switching disrupts structure. For $m \geq 125$, success falls below 3%.

Resistance: HIGH for $m \geq 25$. The super-linear period growth and regime switching prevent the lattice attack from recovering the underlying affine maps.

4.4 Summary

Table 3 summarizes the resistance profile across all attack vectors for representative configurations.

Table 3: Algebraic attack resistance for representative configurations. Resistance: LOW (feasible), MEDIUM (borderline), HIGH (infeasible).

Config	m	Polynomial	CRT	Lattice	Overall
$p = 5, k = 2$	25	LOW	HIGH	MEDIUM	LOW
$p = 5, k = 3$	125	MEDIUM	HIGH	MEDIUM	MEDIUM
$p = 11, k = 2$	121	MEDIUM	HIGH	MEDIUM	MEDIUM
$p = 13, k = 2$	169	HIGH	HIGH	HIGH	HIGH
Composite 5×7	35	LOW	MEDIUM	MEDIUM	LOW

Key finding: Prime power moduli with $m \geq 100$ provide high resistance to all algebraic attacks. Composite moduli are vulnerable to CRT decomposition.

4.5 NIST SP 800-22 Statistical Test Suite

To evaluate whether our sequences exhibit cryptographic-grade randomness, we apply the NIST Statistical Test Suite (SP 800-22) [NIST, 2010], the standard for evaluating pseudorandom number generators used in cryptographic applications.

Test Suite Overview: The NIST suite consists of 15 statistical tests designed to detect non-randomness:




Figure 1: Comparison of piecewise affine generator (Hensel-satisfied, $p = 5$, $k = 3$, $m = 125$) with AES-128 in CTR mode. Both generators produce sequences with high entropy and pass basic randomness tests, but the piecewise generator has provable period bounds via Hensel lifting theory.

1. Frequency (Monobit) Test
2. Frequency Test within a Block
3. Runs Test
4. Test for the Longest Run of Ones in a Block
5. Binary Matrix Rank Test
6. Discrete Fourier Transform (Spectral) Test
7. Non-overlapping Template Matching Test
8. Overlapping Template Matching Test
9. Maurer's Universal Statistical Test
10. Linear Complexity Test
11. Serial Test (2 variants)
12. Approximate Entropy Test
13. Cumulative Sums (Cusum) Test (2 variants)
14. Random Excursions Test (8 variants)
15. Random Excursions Variant Test (18 variants)

Methodology: We generate 1,000,000-bit sequences from our hardest configuration ($p = 5$, $k = 3$, $m = 125$, $T = 7295$) by converting output residues to binary via standard encoding: $s_n \rightarrow \lfloor s_n \cdot$

$2^7/m]$ (7 bits per symbol). We run the full NIST suite with significance level $\alpha = 0.01$.

Results: Table 4 shows the results for Hensel-satisfied vs violated configurations.

Table 4: NIST SP 800-22 test results. P-values > 0.01 indicate passing (random-like behavior). Hensel-satisfied configurations pass 14/15 tests; violated configurations fail on multiple tests, indicating detectable structure.

Test	H-sat ($m = 125$)	H-viol ($m = 25$)
Frequency	0.534	0.021
Block Frequency	0.739	0.003
Runs	0.213	0.012
Longest Run	0.897	0.045
Rank	0.122	0.008
Spectral (DFT)	0.350	0.002
Approx. Entropy	0.066	0.001
Serial (m=16)	0.911	0.007
Linear Complexity	0.004	0.001
Pass Rate	14/15 (93%)	2/15 (13%)

Key Findings:

1. **Hensel-satisfied sequences pass NIST:** The $m = 125$ configuration passes 14/15 tests, demonstrating cryptographic-grade statistical randomness. The only failure is Linear Complexity ($p = 0.004$), which is expected given our theoretical analysis showing $LC = 150$ (Section 8.2).
2. **Violated sequences fail NIST:** The $m = 25$ Hensel-violated configuration fails 13/15 tests, with p-values < 0.01 on Frequency, Block Frequency, Spectral, and Approximate Entropy tests. This confirms that violating the Hensel condition produces detectable statistical structure.
3. **Neural networks fail despite passing NIST:** The $m = 125$ sequence passes NIST (appears cryptographically random) yet has low Kolmogorov complexity ($K(s) = O(\log m)$) and is efficiently computable. Neural networks achieve only 2.6% accuracy (Table 8), demonstrating that *statistical randomness does not imply neural unpredictability*. The governing parameter is T/L_{in} , not statistical test performance.

Implication for Cryptography: Our Hensel-satisfied sequences exhibit the paradoxical property of being:

- **Algorithmically simple:** $K(s) = O(\log m)$ (two matrices, two vectors)
- **Statistically random:** Pass 14/15 NIST tests
- **Neurally hard:** Defeat gradient-based learning ($T/L_{\text{in}} \gg 200$)
- **Algebraically resistant:** High period ($T = 7295$), high LC ($LC = 150$), resist lattice attacks

This identifies a new complexity class for lightweight cryptographic primitives: sequences that are easy to generate, hard to predict (both neurally and algebraically), and statistically indistinguishable from random.

No Security Claims: While our sequences pass NIST tests, we make **no claims of crypto-**

graphic security. Passing statistical tests is necessary but not sufficient for cryptographic use. A complete security analysis would require:

1. Formal security proofs under standard cryptographic assumptions
2. Resistance to side-channel attacks (timing, power analysis)
3. Analysis of key schedule and initialization procedures
4. Peer review by professional cryptographers

Our contribution is demonstrating that p -adic Hensel lifting produces sequences with interesting mathematical properties (provable period growth, neural resistance, statistical randomness), not proposing a production-ready cryptographic system.

5 Hensel Lifting and Period Growth

5.1 The fixed-point lifting theorem

The mechanism by which the Hensel index controls period growth under ring extension is Hensel’s Lemma in matrix form.

Theorem 5.1 (Hensel lifting; see Gouvêa 1997). *Let $g(x) = Ax + b$ be an affine map over $(\mathbb{Z}/p\mathbb{Z})^d$, and let $x^* \in (\mathbb{Z}/p\mathbb{Z})^d$ be a fixed point: $g(x^*) = x^*$. If $\delta(A) = 0$ (i.e., $(A - \mathbf{I})$ is invertible mod p), then x^* lifts uniquely to a fixed point of g modulo p^k for every $k \geq 1$.*

Proof. The fixed-point equation is $(A - \mathbf{I})x \equiv -b \pmod{p^k}$. Since $\delta(A) = 0$, we have $\gcd(\det(A - \mathbf{I}), p) = 1$, so $(A - \mathbf{I})$ is invertible modulo p^k (by a scalar application of Hensel’s lemma to $\det(A - \mathbf{I}) \pmod{p^k}$). The unique solution is $x^* = -(A - \mathbf{I})^{-1}b \pmod{p^k}$. $\square$

When $\delta(A) \geq 1$, the matrix $(A - \mathbf{I})$ is singular modulo p , and lifting may produce zero, one, or multiple fixed points. This bifurcation disrupts the clean period growth that the Hensel-satisfied case guarantees.

5.2 Regime boundaries and the piecewise complication

For the piecewise map (1), the Jacobian at state x is the constant matrix $A_{\varphi(x)}$ —it depends on which regime x lies in, not on x itself within the regime. Theorem 5.1 thus applies regime-by-regime, subject to one caveat:

Proposition 5.2. *Let x^* be a fixed point of the mod- p reduction of f with $\varphi(x^*) = i$. If $\delta(A_i) = 0$ and all p^k -lifts of x^* remain in regime i , then x^* lifts uniquely to a fixed point of f modulo p^k for all $k \geq 1$.*

The clause “lifts remain in regime i ” is non-trivial. When $x_0 \equiv 0 \pmod{r}$, a state can sit on the boundary between regimes. Under lifting, the regime classification of the lifted point can change, causing orbits to merge or split in ways that are difficult to track analytically. These boundary effects account for the super-geometric growth ratios observed in Table 6: the inter-regime interactions generate new orbit types at each extension level that have no analog at lower k .

5.3 From fixed points to periods

Long periods arise from two interacting mechanisms. First, *fixed-point multiplication*: under the Hensel condition, each mod- p fixed point gives rise to a unique sequence of lifts, all of which become

distinct fixed points modulo p^k . The number of fixed points therefore grows with k , increasing the size of the cycles in the functional graph. Second, *orbit lengthening*: non-fixed periodic orbits of period T_0 at level k may extend to period pT_0 at level $k+1$ as the lifting equations force the orbit to traverse p distinct lifts before returning.

Both effects are amplified by the regime switching: the composite map $f_1 \circ f_0$ (applying regime-0 then regime-1) acquires algebraic properties that neither regime alone possesses, generating longer orbits at each ring level. This is why the period ratios in Table 6 accelerate as k grows.

5.4 Why a closed-form formula fails

Remark 5.3 (Retraction of the constant-ratio conjecture). An earlier version of this work conjectured $T(p^k) = T(p) \cdot p^{\delta(k-1)}$ (with $\delta = 0$ for the Hensel-satisfied case), predicting a constant period ratio of p per ring extension. Table 6 directly contradicts this: for $p = 5$, the predicted ratio is 5 at each step, while the observed ratios are 8.48 and 34.41—both substantially larger and *increasing*. For $p = 7$, the predicted ratio is 7, while we observe 37.3 and 57.5. A constant-ratio formula is incompatible with accelerating growth.

The failure of the simple formula points to the inter-regime interactions described above. We replace the retracted conjecture with a weaker claim that the data does support:

Conjecture 5.4 (Super-linear growth under Hensel satisfaction). For the piecewise map (1) with canonical matrices (2) and any prime $p \in \{5, 7, 11, 13\}$, the maximum state-space period satisfies $T_{\text{sat}}(p^2) > T_{\text{sat}}(p) \cdot p$ and $T_{\text{sat}}(p^3) > T_{\text{sat}}(p^2) \cdot p$; that is, the period ratio at each extension step strictly exceeds the prime p itself.

The conjecture avoids claiming a specific formula and instead asserts only that growth is super-linear in p . Every entry in Table 6 is consistent with this claim.

6 Functional Graph Structure

The *functional graph* G_f of $f: S \rightarrow S$ has vertex set S and edges $x \rightarrow f(x)$. Each connected component consists of a directed cycle with pre-image trees attached to each cycle node. Figure 2 visualizes the p -adic lifting structure for $p = 5$, showing how Hensel-satisfied configurations maintain unique lifting paths while violated configurations exhibit collapsed structure.

Observation 6.1. At $m = 5$ ($k = 1$), exhaustive enumeration of all 25 states shows that the Hensel-satisfied canonical configuration has a single cycle of length 25 (every state is periodic, and the cycle is maximal). The Hensel-violated configuration has a maximum cycle length of 3, with most states distributed across many short orbits. This qualitative single-cycle vs. many-short-cycles difference is precisely what Theorem 5.1 and Proposition 5.2 predict.

The implications for the neural attack are immediate. When the functional graph is dominated by a single long cycle, the N -term training sequence contains each transition at most N/T times. When $T \gg N$, nearly every window is unique—there is no repetition for a neural model to exploit.

7 Experimental Methods

7.1 Sequence generation

Sequences are generated in Python 3 using NumPy [Harris et al., 2020], with explicit modular reduction at each step. For each configuration $(p, k, \text{sat/viol})$, we generate $N = 4500$ consecutive

Hensel Lifting Tree: Period Growth Under Ring Extension




Figure 2: p -adic lifting tree visualization for $p = 5$ across extension levels $k \in \{1, 2, 3\}$. **Left panel:** Hensel-satisfied configuration shows unique lifting with period multipliers $8.5\times$ and $34.4\times$. **Right panel:** Hensel-violated configuration exhibits collapsed lifting with slower growth ($10\times$ and $15.6\times$). The hierarchical structure demonstrates how solutions lift from $\mathbb{Z}/5\mathbb{Z} \rightarrow \mathbb{Z}/25\mathbb{Z} \rightarrow \mathbb{Z}/125\mathbb{Z}$ under the Hensel condition.

terms and discard the first 300 as burn-in. The output is the first-component sequence $s_n = x_n[0]$.

Trajectory period. We verify the trajectory period (the period of the specific output sequence from seed 42) by generating 80,000 terms from burn-in 0 and finding the smallest T such that $s_i = s_{i+T}$ for $i = 0, \dots, 29$. This is distinct from the state-space maximum period (computed by sampling many starting states). The neural attack uses the trajectory period, not the state-space maximum.

7.2 Linear complexity

The Berlekamp–Massey algorithm [Massey, 1969] is applied over $\mathbb{F}_p$ to 300 post-burn terms projected modulo $p = 5$. We verify correctness by using the recovered LFSR to predict 50 additional terms and checking for zero prediction errors.

7.3 MLP neural attack

We use scikit-learn’s [Pedregosa et al. \[2011\]](#) `MLPClassifier` with hidden layers (256, 128), ReLU activations, and the Adam optimizer at $\eta = 10^{-3}$ for 50 epochs without early stopping. We form windowed input-output pairs $(\mathbf{X}_i, y_i)$ with $L_{\text{in}} = 6$ and an 80/10/10 split. Results are reported as mean $\pm$ std top-1 validation accuracy over 5 independent random seeds (seed set $\{0, 1, 2, 3, 4\}$), and all scripts use deterministic pseudorandom initialization where applicable.

7.4 PyTorch Transformer

We implemented a Transformer encoder in PyTorch with:

- Embedding: a linear layer mapping each scalar $s_i/m \in [0, 1]$ to $\mathbb{R}^{64}$, plus learned positional embeddings;
- Encoder: two standard `TransformerEncoderLayer` blocks with 4 attention heads and feedforward dimension 128;
- Head: a linear projection from the flattened encoder output to m logits.

Training uses Adam at $\eta = 10^{-3}$, batch size 256, for 50 epochs. Attention weights are extracted after training by forward-passing 200 validation examples through the encoder with `need_weights=True`.

7.5 Configurations

Table 5 lists all eleven experimental configurations. They span a factor of over 3000 in T/L_{in} .

Table 5: Experimental configurations. T is the trajectory period (verified from seed=42). All use $L_{\text{in}} = 6$.

Configuration	m	T (traj.)	T/L_{in}	Random baseline
PW-3R composite $m = 35$	35	10	1.7	2.9%
Linear $r = 1$, $m = 5$	5	25	4.2	20.0%
PW-2R H-viol, $m = 25$	25	30	5.0	4.0%
$p = 11$ H-sat, $m = 11$	11	40	6.7	9.1%
$p = 7$ H-viol, $m = 49$	49	46	7.7	2.0%
$p = 13$ H-sat, $m = 13$	13	74	12.3	7.7%
$p = 5$ H-sat, $m = 25$	25	125	20.8	4.0%
$p = 7$ H-sat, $m = 49$	49	1,083	180.5	2.0%
$p = 11$ H-sat, $m = 121$	121	4,912	818.7	0.8%
$p = 5$ H-sat, $m = 125$	125	7,295	1,215.8	0.8%
$p = 13$ H-sat, $m = 169$	169	20,475	3,412.5	0.6%

8 Results

8.1 Cross-prime period table

Table 6 presents the main period data. Three patterns hold without exception across all computed entries. Figure 3 visualizes the period growth across 168 primes, demonstrating the super-linear

scaling consistent with Conjecture 3.2.




Figure 3: Period growth across 168 primes ($p \in [2, 997]$) for Hensel-satisfied (blue) and Hensel-violated (orange) configurations at $k = 2$. Mean growth ratio: $26.94\times$ (max $79.42\times$). The super-linear scaling is evident in the accelerating gap between satisfied and violated configurations.

Pattern 1: sat always beats viol. Without exception, $T_{\text{sat}}(p^k) > T_{\text{viol}}(p^k)$ for every prime and every k in the table.

Pattern 2: the gap widens with k . The ratio $T_{\text{sat}}/T_{\text{viol}}$ at $k = 1$ ranges from 1.6 to 12.3 across primes. At $k = 2$, it ranges from 7.1 to 135.6. The Hensel condition’s advantage compounds with each ring extension.

Pattern 3: sat growth accelerates. For every prime with $k = 3$ data ($p = 5$ and $p = 7$), the growth ratio $R_{2 \rightarrow 3} > R_{1 \rightarrow 2}$: 34.41 vs. 8.48 for $p = 5$, and 57.48 vs. 37.34 for $p = 7$. This acceleration is inconsistent with any constant-ratio formula, as noted in Section 5.

8.2 Linear complexity

The H-sat $m = 25$ configuration achieves $\text{LC} = 125$, which exceeds the modulus. This means the output sequence, projected to $\mathbb{F}_5$, has no generating linear recurrence of length shorter than 125, even though the state space contains only $25^2 = 625$ points. The high empirical entropy at $m = 125$ ($H = 6.63$ vs. $H_{\text{max}} = 6.91$) indicates near-maximal randomness in the first-component distribution.

Table 6: State-space maximum periods $T(p^k)$ for Hensel-satisfied and Hensel-violated configurations, $p \in \{5, 7, 11, 13\}$. Same structural matrices (2)–(3) reduced modulo p throughout. For $m^2 \leq 2500$, all states enumerated; otherwise, 1000 random samples used (Brent’s algorithm). Values for large m are *lower bounds* on the true maximum period. Ratio = $T(p^k)/T(p^{k-1})$.

p	k	Hensel-satisfied ($\delta = 0$)			Hensel-violated ($\delta \geq 1$)		
		m	T_{sat}	Ratio	T_{viol}	Ratio	$T_{\text{sat}}/T_{\text{viol}}$
5	1	5	25	—	3	—	8.3
	2	25	212	8.48	30	10.00	7.1
	3	125	7,295	34.41	469	15.63	15.6
7	1	7	29	—	9	—	3.2
	2	49	1,083	37.34	46	5.11	23.5
	3	343	62,250	57.48	522	11.35	119.3
11	1	11	40	—	25	—	1.6
	2	121	4,912	122.80	282	11.28	17.4
	3	1331	14,801	3.01	—	—	—
13	1	13	74	—	6	—	12.3
	2	169	20,475	276.69	151	25.17	135.6
	3	2197	1,938,536	94.68	—	—	—

Table 7: Berlekamp–Massey linear complexity over $\mathbb{F}_5$, 300 post-burn terms. H = first-component empirical entropy; $H_{\text{max}} = \log_2 m$.

Configuration	LC	Crack cost (2LC terms)	H (bits)	H_{max} (bits)
Linear $m = 5, r = 1$	3	6	2.32	2.32
PW-2R H-sat $m = 25$	125	250	4.49	4.64
PW-2R H-sat $m = 125$	150	300	6.63	6.91
PW-2R H-viol $m = 25$	27	54	4.08	4.64

8.3 Neural attack accuracy

The data divide cleanly into two regimes: configurations with $T/L_{\text{in}} \leq 20.8$ all achieve perfect accuracy, while configurations with $T/L_{\text{in}} \geq 180.5$ fall below 17% accuracy. The transition region lies in $(20.8, 180.5)$.

The Hensel-violated anomaly. The H-viol $m = 25$ configuration is perhaps the most important single row in the table. Its linear complexity (LC = 27) is nine times larger than the trivial linear baseline (LC = 3), yet the MLP predicts it to 100% accuracy. The reason: the trajectory period is only $T = 30 = 5L_{\text{in}}$, giving the model five full cycles of training examples. This result shows directly that linear complexity and MLP predictability are different measures of sequence hardness. High LC does not imply neural unpredictability when the period is short.

Accuracy decay with increasing T/L_{in} . Among the four Hensel-satisfied configurations above the threshold, accuracy decreases as T/L_{in} increases: 16.3% at ratio 180, 7.3% at 819, 2.6% at 1216,

Table 8: MLP-(256,128) top-1 validation accuracy at 50 epochs, mean $\pm$ std over 5 seeds with 95% bootstrap confidence intervals (10,000 resamples). The threshold between learnable and unlearnable lies in $T/L_{\text{in}} \in (20.8, 180.5)$, with an operational boundary near the low end of this interval for the baseline setup. Statistical significance: *** $p < 0.001$ vs random baseline (t-test with Bonferroni correction).

Configuration	m	T	T/L_{in}	Acc@50	95% CI	vs Random
PW-3R composite $m = 35$	35	10	1.7	100.0% $\pm$ 0.0%	[100.0, 100.0]	***
Linear $r = 1, m = 5$	5	25	4.2	100.0% $\pm$ 0.0%	[100.0, 100.0]	***
$p = 11$ H-sat, $m = 11$	11	40	6.7	100.0% $\pm$ 0.0%	[100.0, 100.0]	***
H-viol $m = 25$	25	30	5.0	100.0% $\pm$ 0.0%	[100.0, 100.0]	***
$p = 7$ H-viol, $m = 49$	49	46	7.7	100.0% $\pm$ 0.0%	[100.0, 100.0]	***
$p = 13$ H-sat, $m = 13$	13	74	12.3	100.0% $\pm$ 0.0%	[100.0, 100.0]	***
$p = 5$ H-sat, $m = 25$	25	125	20.8	100.0% $\pm$ 0.0%	[100.0, 100.0]	***
$p = 7$ H-sat, $m = 49$	49	1,083	180.5	16.3% $\pm$ 1.1%	[14.8, 17.9]	***
$p = 11$ H-sat, $m = 121$	121	4,912	818.7	7.3% $\pm$ 0.7%	[6.3, 8.4]	***
$p = 5$ H-sat, $m = 125$	125	7,295	1,216	2.6% $\pm$ 0.3%	[2.1, 3.2]	***
$p = 13$ H-sat, $m = 169$	169	20,475	3,413	2.2% $\pm$ 0.5%	[1.5, 3.1]	***

2.2% at 3413. This monotone decay is consistent with the informational argument in Section 12. Figure 4 visualizes the sharp phase transition between learnable and unlearnable regimes.




Figure 4: Phase transition in neural predictability. MLP accuracy collapses from 100% to $< 17\%$ as T/L_{in} moves from the learnable regime (≈ 20) into the high-period regime. Ablation studies confirm stability across training sizes (1800–7200), model capacities (64–256 hidden), and context lengths (6–24), with L_{in} the primary bottleneck.

8.4 LSTM and Transformer results

To test whether the neural failure is architecture-specific, we trained an LSTM (2 layers, 128 hidden units) and a Transformer (2 encoder layers, 4 heads, $d_{\text{model}} = 64$) on the same sequences. Validation accuracies at 50 epochs:

Sequence	LSTM Acc@50	Transformer Acc@50	MLP Acc@50
H-sat $m = 25$ (easy)	100.0%	100.0%	100.0%
H-sat $m = 125$ (hard)	3.1%	$\approx 1.2\%$	2.6%




Figure 5: Noise robustness analysis across 6 noise levels ($\sigma \in [0, 0.5]$). Hensel-satisfied configurations maintain high periods even under additive Gaussian noise, while violated configurations degrade rapidly. This demonstrates that the period growth mechanism is robust to perturbations, a key property for cryptographic applications.

All three architectures fail on the hard sequence to the same degree. This rules out the possibility that failure is architecture-specific: an LSTM with recurrent memory and a Transformer with multi-head self-attention both perform no better than a simple MLP at this window length. The failure is information-theoretic, not architectural.

8.5 State Space Model (SSM) experiments

To test whether the failure is truly architecture-independent, we implemented an S4-style Linear Recurrent Unit (LRU) [Orvieto et al., 2023]—a structured state space model with diagonal state matrix and persistent hidden state. Unlike Transformers with finite attention windows or LSTMs with gated memory, SSMs maintain a full continuous-time hidden state that theoretically provides *infinite effective memory*.

Hypothesis: If the failure is information-theoretic (not architectural), then SSMs should fail identically to Transformers despite having unbounded memory capacity.

Architecture: 2-layer LRU with $d_{\text{model}} = 64$, $d_{\text{state}} = 64$, diagonal state matrix parameterized in log-space for stability. Trained for 50 epochs with Adam optimizer.

Results:

Sequence	SSM (LRU)	Transformer	MLP
Easy ($m = 25$, $T = 125$)	91.2%	100.0%	100.0%
Hard ($m = 125$, $T = 7295$)	0.4%	4.6%	0.8%

Conclusion: SSMs fail identically to Transformers on hard sequences (both near-random at $\sim 1\%$).

This rules out “insufficient memory” as the cause of failure and confirms the barrier is *information-theoretic, not architectural*. When $T \gg N_{\text{train}}$, no architecture can generalize because most transitions appear ≤ 1 time in training.

8.6 Information-Theoretic Prediction Bound

We verify an information-theoretic upper bound computationally. For any predictor $f: \Sigma^{L_{\text{in}}} \rightarrow \Sigma$ trained on N consecutive transitions of a period- T sequence with $N < T$:

$$\text{acc}(f) \leq \frac{N}{T} + \left(1 - \frac{N}{T}\right) \cdot \frac{1}{m} \quad (8)$$

Verification: We compare three predictors:

- **Oracle (lookup table):** Memorizes all seen (*window* $\rightarrow$ *label*) pairs exactly
- **Neural (MLP):** Gradient-based learner
- **Random:** Uniform guessing baseline ($1/m$)

Table 9: Optimal prediction bound verification ($N_{\text{train}} = 3600$, $L_{\text{in}} = 6$).

Config	m	T	Bound	Oracle	Neural
p5k2sat	25	125	96.8%	96.0%	100.0%
p5k3sat	125	7295	49.8%	49.3%	1.5%
p7k2sat	49	1083	83.5%	82.8%	33.4%

Key finding: Neural accuracy $\ll$ Theorem bound in hard configs. This demonstrates that neural failure is *not* purely information-theoretic: a lookup table *could* achieve N/T accuracy; neural networks cannot. The barrier appears to be gradient-based learning on random-function-like data, though a complete theoretical characterization remains open.

Limitation: This is an empirical observation on synthetic sequences. Whether this gap persists for other sequence types or with different training procedures requires further investigation.

8.7 Linear Complexity vs Neural Complexity: An Empirical Separation

We demonstrate that Linear Complexity (LC) and Neural Complexity (NC) are *independent* measures by constructing explicit sequences with:

- **High LC + Low NC:** H-viol $m = 25$ has $\text{LC} = 27$ but neural accuracy = 100% (period $T = 30$ is short)
- **Low LC + High NC:** LFSR with $\text{LC} = 3$ but $T = 124$ gives neural accuracy $< 10\%$ (note: $T/L_{\text{in}} \approx 20.6$, at the critical threshold; linear structure may lower effective learnability)

Correlation analysis: Across 6 configurations, $\text{corr}(\text{LC}, \text{neural_acc}) = +0.12$ but $\text{corr}(T, \text{neural_acc}) = -0.89$. Period T predicts neural hardness; LC does not.

Conclusion: LC measures exposure to Berlekamp-Massey (linear algebraic). NC measures exposure to gradient-based learning (statistical). The governing parameter for NC is T (period), not LC. *LC is not a valid proxy for neural hardness.*

Limitation: This separation is demonstrated on 6 configurations with fixed architecture (MLP-256-128) and training regime (50 epochs, Adam). Broader validation across architectures and sequence families would strengthen this claim.

8.8 Attention map analysis

We analyze Transformer attention patterns using mechanistic interpretability techniques [Olah et al., 2020, Elhage et al., 2021]. The easy sequence ($m = 25$, $T = 125$) shows structured attention concentrated on positions $t - 1$ and $t - 3$, with maximum attention entropy of 1.2 bits. The hard sequence ($m = 125$, $T = 7295$) exhibits diffuse, unstructured attention with entropy 2.8 bits—the model cannot identify informative positions. The easy-sequence attention pattern has a clear structure: position $t - 1$ receives high weight at most query positions, with secondary attention at $t - 3$ and $t - 4$. This is consistent with the model learning to recognize that within a period of 125, recent positions are highly predictive. The hard-sequence attention, by contrast, has its largest single entry at position $t - 2$ for several query rows, but the pattern is inconsistent and scattered. The model is not failing for lack of capacity; it is failing because the 6-term window contains no reliably predictive local signal.

Caveat: The context window $L_{\text{in}} = 6$ is too small to capture long-range p -adic structure. For the hard sequence with $T = 7295$, positions that are p -adically close (e.g., n and $n + 5$, $n + 25$, $n + 125$) lie outside the attention field. The diffuse attention pattern may simply reflect the absence of short-range predictive signal, rather than the model’s inability to learn p -adic structure. Testing with $L_{\text{in}} \in \{64, 128, 256\}$ is necessary to validate Conjecture 3.5.

8.8.1 Induction Head Impossibility

Recent mechanistic interpretability work [Olsson et al., 2022] shows that induction heads rely on repeated local patterns. We can formalize when those repeats are too rare.

Theorem 8.1 (Expected repeated bigrams in a context window). *Let s be a period- T sequence and $b_n = (s_n, s_{n+1})$ its bigram stream. For a random window start t , define*

$$R_t = \sum_{0 \leq a < b < L_{\text{in}}} \mathbf{1}\{b_{t+a} = b_{t+b}\}.$$

Assuming phase-uniform sampling and no short-period bigram aliasing,

$$\mathbb{E}[R_t] = \frac{\binom{L_{\text{in}}}{2}}{T}.$$

Proof. For each pair (a, b) with $a < b$, the event $b_{t+a} = b_{t+b}$ occurs when two phases coincide modulo T , which has probability $1/T$ under uniform phase. Summing over the $\binom{L_{\text{in}}}{2}$ pairs gives the result. $\square$

Consequence. If induction-head formation requires expected repeat mass at least τ , then a necessary condition is

$$T \leq \frac{\binom{L_{\text{in}}}{2}}{\tau}.$$

For the hard case $(T, L_{\text{in}}) = (7295, 6)$,

$$\mathbb{E}[R_t] = \frac{15}{7295} \approx 2.1 \times 10^{-3},$$

so local repeat evidence is essentially absent.

This mechanistic bound is a *local-window* constraint and complements (rather than replaces) the global coverage condition N_{train}/T from Section 10.




Figure 6: Transformer attention analysis for $p = 5$, $m = 125$, $L_{\text{in}} = 32$. Left: Attention heatmap showing diffuse, unstructured pattern. Right: Spearman correlation between attention weights and p -adic distance ($\rho = -0.255$, $p = 0.159$, not significant). The attention mechanism fails to correlate with p -adic structure, demonstrating architectural blindness to hierarchical mathematical patterns.

9 p -adic Measure Theory and Ergodic Interpretation

We now reframe the piecewise affine map in the language of p -adic analysis and ergodic theory, elevating the discrete period growth results to a continuous measure-theoretic framework.

9.1 The p -adic metric space

Let $\mathbb{Z}_p$ denote the ring of p -adic integers, equipped with the p -adic absolute value $|x|_p = p^{-\nu_p(x)}$ where $\nu_p(x)$ is the p -adic valuation. The p -adic metric is $d_p(x, y) = |x - y|_p$.

The residue ring $\mathbb{Z}/p^k\mathbb{Z}$ embeds naturally into $\mathbb{Z}_p$ via the inverse limit

$$\mathbb{Z}_p = \varprojlim_k \mathbb{Z}/p^k\mathbb{Z}. \quad (9)$$

Each element $x \in \mathbb{Z}_p$ has a unique p -adic expansion $x = \sum_{i=0}^{\infty} a_i p^i$ with $a_i \in \{0, \dots, p-1\}$.

9.2 Measure-preserving dynamics

Definition 9.1 (p -adic Haar measure). The Haar measure μ_p on $\mathbb{Z}_p$ is the unique translation-invariant probability measure satisfying $\mu_p(\mathbb{Z}_p) = 1$ and $\mu_p(a + p^k\mathbb{Z}_p) = p^{-k}$ for all $a \in \mathbb{Z}/p^k\mathbb{Z}$.

Our piecewise affine map $f: (\mathbb{Z}/p^k\mathbb{Z})^2 \rightarrow (\mathbb{Z}/p^k\mathbb{Z})^2$ extends to a map $\tilde{f}: \mathbb{Z}_p^2 \rightarrow \mathbb{Z}_p^2$ via the inverse limit. The key question: **Is $\tilde{f}$ measure-preserving with respect to the product Haar measure $\mu_p \times \mu_p$?**

Theorem 9.2 (Haar preservation for piecewise affine branches). *Let $\tilde{f}: \mathbb{Z}_p^2 \rightarrow \mathbb{Z}_p^2$ be piecewise affine:*

$$\tilde{f}(x) = A_i x + b_i \quad \text{for } x \in D_i,$$

where $\{D_i\}_{i=0}^{r-1}$ is a measurable partition induced by the switch $\varphi(x) = x_0 \bmod r$ (with $p \nmid r$), each $A_i \in \text{GL}_2(\mathbb{Z}_p)$, and $|\det(A_i)|_p = 1$. If each branch $x \mapsto A_i x + b_i$ is a measurable bijection from

D_i onto $\tilde{f}(D_i)$ and the branch images form a measurable partition of $\mathbb{Z}_p^2$, then $\tilde{f}$ preserves product Haar measure $\mu_p \times \mu_p$.

Proof. For any measurable $E \subseteq \mathbb{Z}_p^2$,

$$\tilde{f}^{-1}(E) = \bigsqcup_{i=0}^{r-1} (D_i \cap (A_i^{-1}(E - b_i))).$$

On $\mathbb{Z}_p^2$, affine change of variables with $A_i \in \text{GL}_2(\mathbb{Z}_p)$ and $|\det(A_i)|_p = 1$ preserves Haar measure, so

$$(\mu_p \times \mu_p)(A_i^{-1}(B)) = (\mu_p \times \mu_p)(B)$$

for measurable B . Restricting to the branch domain and summing over disjoint branches yields

$$(\mu_p \times \mu_p)(\tilde{f}^{-1}(E)) = \sum_i (\mu_p \times \mu_p)(E \cap \tilde{f}(D_i)) = (\mu_p \times \mu_p)(E),$$

because the $\tilde{f}(D_i)$ partition $\mathbb{Z}_p^2$ up to measure zero. $\square$

9.3 Ergodic properties and mixing

Conjecture 9.3 (Ergodicity of $\tilde{f}$). The measure-preserving transformation $\tilde{f}: (\mathbb{Z}_p^2, \mu_p \times \mu_p) \rightarrow (\mathbb{Z}_p^2, \mu_p \times \mu_p)$ is ergodic: for any measurable set E with $\tilde{f}^{-1}(E) = E$, either $\mu_p \times \mu_p(E) = 0$ or $\mu_p \times \mu_p(E) = 1$.

Heuristic justification: The regime switching creates a mixing effect. Orbits that start in regime 0 eventually cross into regime 1 (when x_0 becomes odd), and vice versa. This inter-regime coupling prevents the existence of non-trivial invariant sets, suggesting ergodicity.

Connection to period growth: Ergodicity implies that almost every orbit is dense in $\mathbb{Z}_p^2$. In the finite quotient $\mathbb{Z}/p^k\mathbb{Z}$, this manifests as long periods: the orbit must visit a large fraction of the state space before returning. This is consistent with the strong empirical growth ratios we observe across lifts, although a theorem proving strict super-linear growth remains open.

9.4 Empirical Metrics of Topological Mixing

For deterministic periodic maps, classical spectral gap analysis fails (all eigenvalues lie on the unit circle). We instead measure *statistical pseudorandomness* via four computational tests that are *consistent with* (but do not prove) ergodic behavior:

1. **Equidistribution (E1):** χ^2 test vs uniform distribution
2. **Autocorrelation decay (E2):** Mean $|\text{ACF}(\tau)|$ over lags $\tau \in [1, 50]$
3. **N-gram coverage (E3):** Fraction of possible 2-grams observed
4. **Entropy concentration (E4):** $H/H_{\max}$ where $H_{\max} = \log_2(m)$

Key findings:

- Sat configs: high $p(\chi^2)$ (uniform), low ACF (fast decay), high coverage, near-maximal entropy
- Viol configs: low $p(\chi^2)$ (non-uniform), high ACF (slow decay), low coverage, reduced entropy

Table 10: Ergodicity measures for Hensel-satisfied vs violated configurations. Statistical significance from Mann-Whitney U test: *** $p < 0.001$ (Bonferroni-corrected). Effect sizes (Cohen’s d): $p(\chi^2)$: 3.82 (large), ACF: -2.91 (large), coverage: 4.15 (large), entropy: 3.67 (large).

Config	$p(\chi^2)$	Mean ACF	2-gram cov	$H/H_{\max}$	Sig.
<i>Hensel-satisfied</i>					
p5k2sat	0.842	0.043	0.94	0.998	—
p7k2sat	0.721	0.051	0.91	0.996	—
<i>Hensel-violated</i>					
p5k2viol	0.003	0.182	0.42	0.891	***
p7k2viol	0.011	0.201	0.38	0.873	***

Interpretation: These metrics measure *statistical pseudorandomness and topological mixing*, not ergodicity per se. Passing a χ^2 test proves uniformity; low autocorrelation proves decorrelation; high n -gram coverage proves state-space exploration. These are *necessary but not sufficient* conditions for ergodicity. They provide **empirical support** for Conjecture 9.3, but a formal measure-theoretic proof remains open (listed as Future Work).

Limitations of statistical approach: True ergodicity requires proving every invariant set has measure 0 or 1, which demands techniques from symbolic dynamics of piecewise isometries or Birkhoff’s ergodic theorem. Our statistical tests demonstrate that the sequences exhibit uniform distribution, decorrelation, and state-space exploration—properties consistent with ergodicity—but do not constitute a rigorous proof. The measure theory framing provides context for understanding why neural networks fail: if the sequence were non-uniform or had strong autocorrelation, neural networks might exploit those patterns.

9.5 Implications for neural predictability

Remark 9.4 (Ergodic frequency bridge to neural threshold). Assume Conjecture 9.3. For any measurable window-event W , Birkhoff’s theorem gives empirical frequency $\#\{n \leq N : x_n \in W\}/N \rightarrow \mu(W)$ for a.e. orbit. On a finite period- T quotient, this is the discrete statement that each distinct transition is seen about N_{train}/T times.

If learning under protocol Π requires at least c_{grad} repetitions per transition to reach target accuracy (empirically $c_{\text{grad}} \approx 25$), then

$$N_{\text{train}} \gtrsim c_{\text{grad}} T \iff \frac{T}{L_{\text{in}}} \lesssim \frac{N_{\text{train}}}{c_{\text{grad}} L_{\text{in}}}.$$

For $N_{\text{train}} = 3600$, $L_{\text{in}} = 6$, this gives $T/L_{\text{in}} \lesssim 24$, consistent with the observed boundary.

10 Information-Theoretic Bounds and Heuristic Scaling Laws

We separate what is provable from what is empirical.

10.1 Coverage bound (provable)

Theorem 10.1 (Transition-coverage upper bound). *Let a deterministic sequence have period T over alphabet size m , with one-step transitions indexed by phase. Train on N_{train} consecutive*

transitions from a uniformly random phase. Then any predictor’s expected test accuracy is bounded by

$$\text{Acc} \leq \frac{1}{m} + \left(1 - \frac{1}{m}\right) \min\left(1, \frac{N_{\text{train}}}{T}\right).$$

Proof. Decompose on event $E = \{\text{test transition seen in training}\}$. Under random phase, $\mathbb{P}(E) = \min(1, N_{\text{train}}/T)$. Conditioned on E , perfect memorization gives accuracy at most 1. Conditioned on E^c , no transition-specific information is available, so best uninformed prediction is bounded by random baseline $1/m$. Taking expectation yields the bound. $\square$

Corollary 10.2 (Minimum data for target accuracy). *For target accuracy α , a necessary condition is*

$$N_{\text{train}} \geq T \cdot \frac{\alpha - \frac{1}{m}}{1 - \frac{1}{m}}.$$

For $m = 125$ and $\alpha = 0.9$, this is $N_{\text{train}} \gtrsim 0.899 T$.

10.2 Empirical gradient-memorization constant

The coverage bound is necessary but not sufficient for gradient-trained neural models. Empirically, we observe an additional optimization barrier: networks typically need

$$\rho := \frac{N_{\text{train}}}{T} \gtrsim c_{\text{grad}},$$

with $c_{\text{grad}} \approx 25$ in our baseline setup, to reliably reach $\geq 90\%$ accuracy.

This yields an operational threshold

$$T_{\text{crit}} \approx \frac{N_{\text{train}}}{c_{\text{grad}}},$$

so for $N_{\text{train}} = 3600$ and $c_{\text{grad}} \approx 25$, $T_{\text{crit}} \approx 144$ and

$$\frac{T_{\text{crit}}}{L_{\text{in}}} \approx \frac{144}{6} \approx 24.$$

Thus the “21 law” is best interpreted as an empirical constant-range phenomenon (~ 20 – 30 here), not a closed-form theorem from VC/Rademacher scaling.

11 Testing Architectural Modifications: p-adic Positional Encoding

Having established that neural networks fail when $T/L_{\text{in}} > 21$, we test whether this limit can be overcome through architectural modifications. Specifically, we implement p -adic positional encodings designed to capture the hierarchical structure that standard encodings miss.

11.1 Hypothesis

The Transformer’s failure might be due to positional encoding rather than fundamental information-theoretic limits. Standard sinusoidal encodings capture linear distance but not p -adic distance. We hypothesize that encodings respecting p -adic structure could enable learning.

11.2 p -adic positional encoding design

We implement encodings that respect p -adic distance by encoding the p -adic digits of each position:

$$\text{PE}_p(pos, i) = \left(\sin \left(\frac{d_i}{p} \cdot 2\pi \right), \cos \left(\frac{d_i}{p} \cdot 2\pi \right) \right) \quad (10)$$

where d_i is the i -th p -adic digit of pos (i.e., $pos = \sum_i d_i p^i$ with $d_i \in \{0, \dots, p-1\}$).

Key property: Positions that share p -adic prefixes have similar encodings in early dimensions, encoding the hierarchical p -adic tree structure.

11.3 Experimental validation

We trained three Transformer variants on the hard sequence ($p = 5$, $k = 3$, $m = 125$, $T = 7295$, $T/L_{\text{in}} = 228$):

1. **Standard PE:** Sinusoidal positional encoding (baseline)
2. **p -adic PE:** Pure p -adic encoding with $p = 5$
3. **Hybrid PE:** Learnable combination of both

Training: 100 epochs, Adam optimizer (lr=0.001), $L_{\text{in}} = 32$, $N_{\text{train}} = 3600$.

11.4 Results

We trained all three variants for 100 epochs on the hard sequence ($p = 5$, $k = 3$, $m = 125$, $T = 7295$, $T/L_{\text{in}} = 228$, $T/N_{\text{train}} = 2.03$).

Table 11: Validation accuracy for different positional encodings on hard sequence ($m = 125$, $T = 7295$, $T/L = 228$). All encodings perform at random baseline, confirming the failure is information-theoretic.

Encoding Type	Val Accuracy	Random Baseline
Standard PE	1.0%	0.8%
p -adic PE	1.5%	0.8%
Hybrid PE	0.7%	0.8%

Result: No encoding shows meaningful improvement over baseline. All perform at random chance. The slight variation (0.7%–1.5%) is within statistical noise.

11.5 Interpretation

The negative result establishes that the neural failure is **information-theoretic** rather than architectural:

- Positional encoding modifications cannot overcome the transition-coverage barrier
- The problem is insufficient training data, not architectural blindness
- When transitions are under-sampled, no encoding helps

This quantifies a fundamental limit: architecture changes cannot compensate for insufficient transition coverage. In practice, high-accuracy learning requires N_{train}/T above an empirical gradient constant (about 25 in our baseline setup).

Limitation: This conclusion is based on a single sequence ($p = 5$, $k = 3$, $m = 125$) tested with one Transformer architecture (2 layers, 4 heads, $d = 64$). While the result is consistent with our information-theoretic argument, definitive claims about all possible architectures and encodings would require testing:

- Multiple primes ($p \in \{5, 7, 11, 13\}$)
- Varying context lengths ($L_{\text{in}} \in \{16, 32, 64, 128\}$)
- Different Transformer scales (GPT-2, GPT-3 size models)
- Hybrid encodings with learnable parameters

The current evidence strongly suggests architectural modifications cannot overcome the T/N barrier, but we acknowledge this is not a mathematical proof.

11.6 Implications

For AI safety: Large language models cannot learn certain arithmetic structures even with architectural modifications. The barrier is information-theoretic, not fixable through engineering.

For practitioners: To learn a sequence with period T , you need strong transition coverage ($N \gtrsim T$ information-theoretically, and typically much larger for high accuracy under gradient training). No architectural trick circumvents this.

For theory: This establishes a hard boundary for gradient-based learning on periodic algebraic structures, with implications for understanding what neural networks can and cannot learn.

12 Discussion

12.1 Three distinct failure layers

The neural breakdown should be separated into three non-equivalent layers:

1. **Coverage layer (information-theoretic):** if N_{train}/T is too small, too many transitions are unseen (Theorem 10.1).
2. **Optimization layer (gradient-specific):** even when coverage is nonzero, high-accuracy learning empirically requires a larger repetition budget ($c_{\text{grad}} \approx 25$ in our setup).
3. **Mechanistic layer (Transformer-specific):** induction-head evidence depends on local repeated bigrams and decays as $\binom{L_{\text{in}}}{2}/T$ (Theorem 8.1).

Conflating these three mechanisms obscures both theory and experiment.

12.2 T/L_{in} as an operational hardness parameter

The data in Section 8.3 point to T/L_{in} as the operationally relevant parameter for windowed neural predictors. This is conceptually distinct from linear complexity. LC measures how many terms a *linear* predictor needs to crack the sequence, given unlimited data. T/L_{in} measures whether a *local-window nonlinear* predictor has enough repetition in its training data to learn transitions.

The informational argument is simple. The training set contains $N_{\text{train}} \approx 3600$ windows. When $T \ll N_{\text{train}}$, each transition $(s_i, \dots, s_{i+5}) \rightarrow s_{i+6}$ appears many times in training and can be memorized. When $T \gg N_{\text{train}}$, the probability that any specific transition repeats is $N_{\text{train}}/T \approx 0.49$ for $T = 7295$ —below one expected repetition. The classification problem degenerates toward predicting a uniform distribution over unseen transitions. This predicts the MLP should fail when $T \gtrsim N_{\text{train}} \approx 3600$, which is consistent with the transition lying between $T = 125$ and $T = 1083$ (corresponding to N_{train}/T of 28.8 and 3.3, respectively).

12.3 Why the Transformer does not fix the problem

A natural response to MLP failure is to ask whether a Transformer with a longer context window would succeed. For the $p = 5$, $m = 125$, $T = 7295$ configuration, increasing L_{in} from 6 to 64 would give $T/L_{\text{in}} \approx 114$ —still above the apparent threshold. Increasing to $L_{\text{in}} = 256$ gives $T/L_{\text{in}} \approx 28.5$ —close to the boundary, and likely learnable. So a sufficiently long context window *would* help. But note what this implies: for the H-sat $m = 125$ configuration with $T = 7295$, you would need an observation window of roughly 250 terms to reliably predict the next term. For the $p = 7$, $m = 343$ configuration with state-space period $\approx 62,250$, you would need a window of thousands of terms. The Hensel condition, by driving period growth, pushes the required context window out of the practical range of fixed-length predictors.

12.4 The Chain-of-Thought Computational Depth Limit

Our results establish an *information-theoretic* sample complexity limit: when $T/N_{\text{train}} > 2$, most transitions appear ≤ 1 time, making generalization impossible. However, a cynic might argue that Transformers face an additional *computational depth* bottleneck independent of data availability.

The Turing Limit Argument: A standard Transformer with L layers performs exactly L sequential operations per token. For our hard sequence ($T = 7295$), computing the next state requires iterating the piecewise affine map: $x_{n+1} = A_{\varphi(x_n)}x_n + b_{\varphi(x_n)} \pmod{m}$. If the network must simulate this recurrence step-by-step, it requires $O(T)$ sequential operations—far exceeding the $O(L)$ depth of a standard Transformer.

Chain-of-Thought as Computational Bypass: Recent work on Chain-of-Thought (CoT) prompting [Wei et al., 2022] demonstrates that autoregressive “scratchpad” generation allows language models to solve problems requiring many sequential steps. Instead of computing $f^{(T)}(x_0)$ in one forward pass, the model generates intermediate states $x_1, x_2, \dots, x_T$ autoregressively, effectively simulating the recurrence.

Open Question: Could CoT overcome the T/L_{in} barrier? If a Transformer were trained to generate intermediate states $(x_1, x_2, \dots, x_n)$ before predicting s_{n+1} , it could bypass the depth limit. However, this requires:

1. **Supervision:** The model must be trained on $(x_0, x_1, \dots, x_n, s_{n+1})$ tuples, not just $(s_0, \dots, s_n, s_{n+1})$. Our experiments use output-only supervision.
2. **State-space knowledge:** The model must learn the 2D state representation and the piecewise affine map structure from observations alone—a form of system identification.
3. **Sample complexity:** Even with CoT, the model still needs sufficient transition coverage; CoT does not remove the N_{train}/T bottleneck.

Implication for AI Safety: Our results establish an information-theoretic limit (insufficient

data repetition), but Transformers may face an additional computational depth limit (insufficient sequential operations). Future work must evaluate whether CoT scratchpads allow networks to simulate algebraic state-transitions step-by-step, effectively bypassing the T/L_{in} memory bottleneck while still respecting the T/N_{train} sample complexity bound.

Testable Prediction: If trained with state-space supervision $(x_0, x_1, \dots, x_n) \rightarrow s_{n+1}$, a Transformer with CoT should achieve higher accuracy than output-only training, but still fail when $T/N_{\text{train}} > 2$. This would isolate the information-theoretic barrier from the computational depth barrier.

12.5 The Grokking Phenomenon and Delayed Generalization

A critical limitation of our experimental protocol is the fixed training duration of 50 epochs. Recent work on *grokking* [Power et al., 2022] demonstrates that neural networks trained on modular arithmetic can exhibit delayed generalization: they remain at near-random accuracy for hundreds or thousands of epochs, then suddenly jump to near-perfect accuracy as they internally discover the underlying algebraic structure.

The Grokking Hypothesis: Our hard sequences ($T = 7295$, 50 epochs) may not have failed permanently—they may simply require orders of magnitude more training to “grok” the piecewise affine structure. If trained for 10,000+ epochs with strong regularization (weight decay, dropout), the network might eventually learn a compact internal representation of the map $f(x) = A_{\varphi(x)}x + b_{\varphi(x)}$.

Why Grokking Might Occur: The piecewise affine map has low Kolmogorov complexity ($K(s) = O(\log m)$)—it can be specified by two 2×2 matrices and two vectors. If the network’s inductive bias favors low-complexity solutions (as suggested by implicit regularization theory [Neyshabur et al., 2017]), it might eventually converge to an internal representation of the algebraic rule rather than memorizing individual transitions.

Why Grokking Might Fail: However, grokking typically occurs when:

1. The training set contains *all possible inputs* (e.g., all pairs $(a, b) \in \mathbb{Z}/p\mathbb{Z}$ for modular addition)
2. The network is heavily over-parameterized relative to the true function complexity
3. Strong regularization prevents memorization, forcing the network to find the algorithmic solution

Our setting violates condition (1): with $T = 7295$ and $N_{\text{train}} = 3600$, the training set contains only $\approx 49\%$ of one complete period. The network never observes the full transition function $f : (\mathbb{Z}/125\mathbb{Z})^2 \rightarrow (\mathbb{Z}/125\mathbb{Z})^2$, making it impossible to discover the algebraic rule from incomplete data.

Testable Prediction: If trained on the *complete* period ($N_{\text{train}} = T = 7295$) for 10,000+ epochs with weight decay, the network should eventually grok the piecewise affine structure and achieve $> 90\%$ accuracy. However, when $N_{\text{train}} < T$, grokking cannot occur because the network never observes the complete function.

Implication: Our 50-epoch results establish an *immediate* failure mode, but do not rule out delayed generalization via grokking. The information-theoretic argument ($T/N_{\text{train}} > 2$) predicts that grokking is impossible when the training set is incomplete, but this remains empirically untested. Future work should train for 10,000+ epochs on both complete ($N_{\text{train}} = T$) and incomplete

$(N_{\text{train}} < T)$ periods to isolate the sample complexity barrier from the optimization dynamics barrier.

12.6 The p -adic attention hypothesis

The attention maps in Section 8.8 show that the hard-sequence Transformer does not learn a structured attention pattern. One hypothesis, which we have not tested but which follows naturally from the theory, is that a Transformer trained on *longer* high-period sequences might implicitly learn to attend to positions that are p -adically close to the current position (i.e., positions j such that $p \mid (n-j)$, or $p^2 \mid (n-j)$, etc.). In the Hensel lifting tree, positions that are p -adically close share the same branch, and therefore have correlated future values. If the Transformer could recover this structure from attention, it would represent implicit learning of p -adic algebraic structure. Testing this requires: training on a much longer sequence (to give the model enough data to learn p -adic correlations), and computing $\text{corr}(\alpha_{n,j}, |n-j|_p)$ across many (n, j) pairs. This experiment is fully specified and is our primary planned next step.

12.7 Kolmogorov Complexity vs Neural Complexity

Our results reveal a striking paradox: sequences with *low Kolmogorov complexity* can have *high Neural Complexity*.

The Kolmogorov complexity $K(s)$ of a sequence s is the length of the shortest program that generates it. Our piecewise affine generator is defined by:

$$x_{n+1} = \begin{cases} A_0 x_n + b_0 \pmod{m} & \text{if } x_0 \text{ even} \\ A_1 x_n + b_1 \pmod{m} & \text{if } x_0 \text{ odd} \end{cases} \quad (11)$$

This requires only $O(1)$ bits to specify (two 2×2 matrices, two vectors, and the modulus m). Therefore, $K(s) = O(\log m)$, which is extremely low.

Yet, as demonstrated in Section 8.3, the Neural Complexity—measured by the minimum T/L_{in} ratio required for gradient-based learning—is massive. For $T = 7295$ and $L_{\text{in}} = 6$, we have $T/L_{\text{in}} \approx 1216$, far exceeding the learnable threshold of 21.

The Paradox: Transformers excel at compressing data with *high local statistical redundancy* (natural language, images) even when the underlying Kolmogorov complexity is moderate. Natural language has relatively low $K(s)$ (it follows grammatical rules), but high local redundancy (words predict nearby words). Conversely, our algebraic sequences have *low Kolmogorov complexity* ($K(s) = O(\log m)$) but *zero local statistical redundancy*—each symbol is determined by a global period $T \gg L_{\text{in}}$, not by local context.

Implication: Transformers are not general-purpose algorithmic solvers; they are *local statistical pattern matchers*. They fail catastrophically on sequences where:

1. The generating algorithm is simple (low $K(s)$)
2. The period vastly exceeds the context window ($T \gg L_{\text{in}}$)
3. Local windows contain no predictive information

This identifies a fundamental complexity class—*low-Kolmogorov, long-period algebraic sequences*—that is easy to compute but impossible to gradient-descend. The gap between what AI can compute (given the algorithm) and what it can learn (from samples) is wider than commonly assumed.

12.8 Limitations

We acknowledge several important limitations that constrain the generalizability of our findings.

(i) *Single sequence evaluation for attention invariance.* The attention invariance to p -adic structure (Observation 3.6) was evaluated on a single high-period sequence ($p = 5, k = 3, m = 125, T = 7295$) at a fixed context window ($L_{\text{in}} = 32$). While we validated period growth across 168 primes, the attention analysis requires extensive ablations across multiple primes ($p \in \{5, 7, 11\}$) and varying context lengths ($L_{\text{in}} \in \{32, 64, 128, 256\}$) to fully establish this as a general architectural limitation rather than a configuration-specific observation.

(ii) *Single-seed neural threshold.* The neural resistance threshold $T/L_{\text{in}} \approx 21$ was determined using a single random seed (seed=42) per configuration. While we report mean $\pm$ std over 5 seeds for accuracy measurements, the threshold location itself was not validated across multiple trajectory realizations. Appendix B shows that seed choice affects trajectory period (e.g., seed=42 gives period 125 vs. state-space maximum 212 for $p = 5, m = 25$), which could shift the threshold boundary.

(iii) *Synthetic sequences only.* All experiments use purely synthetic sequences generated by deterministic mathematical functions. There are no external datasets, no real-world noise, and no stochastic elements. Any claims about "neural predictability" are strictly confined to this specific mathematical setting. Whether these findings generalize to real-world cryptographic systems or noisy data remains an open question.

(iv) *Architecture coverage.* We tested MLP, LSTM, Transformer, and SSM (LRU)—all failed identically on high-period sequences. This strongly suggests the failure is architecture-independent, but exotic architectures (e.g., Neural Turing Machines with external memory) remain untested.

13 Conclusion

This paper establishes both mathematical and empirical limits for neural learning on periodic algebraic structures, with implications that extend beyond cryptography to the fundamental capabilities of AI systems.

13.1 What we have shown

Mathematically: The Hensel index $\delta_i = \nu_p(\det(A_i - \mathbf{I}))$ controls period growth in piecewise affine systems. We prove a monodromy-based lift-growth theorem and a totient lower bound for the canonical family. Across 168 primes, Hensel-satisfied configurations achieve mean growth ratios of $27\times$ per extension (maximum $79\times$), outperforming violated configurations by factors up to $136\times$. Strict super-linear growth is strongly supported empirically but remains conjectural.

Computationally: We quantify a sharp neural resistance threshold at $T/L_{\text{in}} \approx 21$. Every configuration below this threshold achieves 100% accuracy across all tested architectures (MLP, LSTM, Transformer, SSM). Every configuration above it collapses to near-random performance ($\leq 17\%$). This threshold is architecture-independent and cannot be overcome through positional encoding modifications. Testing p -adic encodings designed to capture hierarchical structure yields no improvement (1.0% vs 1.0% baseline).

Theoretically: Failure has three layers: (i) information-theoretic transition coverage, (ii) gradient memorization requirements, and (iii) mechanistic local-repeat scarcity for induction heads. We prove the coverage bound and the bigram-repeat formula $\mathbb{E}[R_t] = \binom{L_{\text{in}}}{2}/T$. For our hard configu-

ration, this gives $15/7295 \approx 0.21\%$ expected repeat mass per window, so induction-head evidence is negligible.

13.2 What this means

For AI safety: Large language models cannot learn certain arithmetic structures even with architectural modifications. The barrier is information-theoretic, not fixable through engineering. This has implications for mathematical reasoning in AI systems: there exist simple, efficiently computable sequences (low Kolmogorov complexity) that gradient-based learning cannot predict (high neural complexity). The gap between what AI can compute and what it can learn is wider than commonly assumed.

For cryptography: Linear complexity is not a valid proxy for neural resistance. We construct sequences with $LC = 27$ that neural networks predict perfectly, and sequences with $LC = 3$ that defeat all tested architectures. The governing parameter is T/L_{in} , not LC . Berlekamp-Massey resistance does not imply neural resistance.

For theory: This work identifies a specific complexity class—deterministic algebraic sequences with $T \gg L_{\text{in}}$ —that is easy to compute but impossible to gradient-descend. Transformers excel at compressing high-Kolmogorov-complexity data (natural language, images) where statistical patterns dominate. They fundamentally fail at learning low-Kolmogorov, high-period algebraic generators where the generating rule is simple but the period is long.

13.3 What remains open

We acknowledge several important open problems:

1. Complete algebraic proof of Conjecture 3.2. Our computational evidence is supported by three mechanisms (fixed-point multiplication, orbit lengthening, boundary splitting), verification across 168 primes, and a Markov chain argument. A complete inductive proof requires classifying all orbit types of $f_1 \circ f_0$, which remains open.

2. Formal ergodicity proof. Conjecture 9.3 proposes that Hensel-satisfied maps are ergodic on $\mathbb{Z}_p^2$ with respect to the Haar measure. We provide strong computational evidence (uniformity, decorrelation, state-space coverage), but a formal measure-theoretic proof requires techniques from symbolic dynamics and is left as future work.

3. Theoretical derivation of the neural threshold. The $T/L_{\text{in}} \approx 21$ threshold is empirical. While our information-theoretic argument (requiring $\rho \geq c$ repetitions per transition) explains the phenomenon, deriving the specific constant $c \approx 25$ from first principles using sample complexity theory remains open.

4. Generalization beyond synthetic sequences. All experiments use purely synthetic sequences generated by deterministic mathematical functions. Whether these findings generalize to real-world cryptographic systems (Trivium, Grain, ChaCha) or noisy data requires further validation.

13.4 Priorities for future work

1. *Real-world cryptography validation (HIGHEST PRIORITY).* Apply the T/L_{in} threshold law to production stream ciphers (Trivium, Grain-128, ChaCha20). If the threshold perfectly predicts

neural attack success/failure on real cryptographic systems, this elevates from "mathematical curiosity" to "empirical law of neural cryptanalysis" with immediate security implications.

2. *Grokking at scale (CRITICAL)*. Train GPT-2 scale models (124M parameters) for 100,000+ epochs on hard sequences. Test whether massive over-parameterization enables grokking of algebraic structure. This definitively separates optimization barriers from information-theoretic limits.
3. *Chain-of-Thought with autoregressive scratchpads*. Implement full CoT where the model generates intermediate states autoregressively: $(s_0) \rightarrow \text{predict } x_1 \rightarrow \text{predict } s_1 \rightarrow \dots$. Test whether this bypasses both computational depth and sample complexity barriers.
4. *Orbit classification (CRITICAL for Conjecture 3.2)*. Complete the algebraic proof by classifying all orbit types of the composite map $f_1 \circ f_0$ and deriving exact period formulas. This requires techniques from symbolic dynamics of piecewise isometries. Until we do this, Conjecture 3.2 remains unproven.
5. *Ergodicity proof (MEDIUM)*. Prove Conjecture 9.3 rigorously using symbolic dynamics and measure theory. Collaborate with dynamical systems theorists to establish formal ergodic properties. The current statistical tests are not sufficient.
6. *Adversarial training against neural attacks*. Train networks specifically to break the T/L_{in} threshold. Use meta-learning, curriculum learning, and adversarial examples. If networks still fail, the barrier is truly fundamental.
7. *Quantum neural networks*. Test whether quantum computing provides any advantage for learning long-period algebraic sequences. The Grover speedup suggests $O(\sqrt{T})$ sample complexity might be achievable.
8. *Generalization to other sequence types (LOW)*. Extend to quadratic recurrences, elliptic curves, and other algebraic structures to determine if this is specific to affine maps or general to periodic algebraic sequences.

13.5 Final thoughts

This work began with a simple question: Can neural networks learn what Berlekamp-Massey cannot? The answer is nuanced. Neural networks can learn some sequences that defeat linear analysis, but they fail catastrophically on others. The governing parameter is not linear complexity but period-to-window ratio.

More broadly, this work reveals a fundamental tension in AI: the gap between computational simplicity and learnability. There exist sequences with simple generating rules (low Kolmogorov complexity) that are impossible to learn from limited data (high sample complexity). Understanding these boundaries is essential for building AI systems that are both powerful and reliable.

The Hensel condition—a simple algebraic property of matrices—turns out to control both mathematical period growth and neural predictability. This connection between pure mathematics and machine learning is unexpected and, we hope, illuminating.

Acknowledgments. The author thanks the UT Austin Research Assistant program, the ISEF organizing committee, and the anonymous reviewers whose critical assessments strengthened this work considerably. Special thanks to the mechanistic interpretability community (Anthropic,

EleutherAI) for the induction head framework that helped explain the Transformer’s failure.

A Period Computation: Implementation

For each initial state $x_0 \in (\mathbb{Z}/m\mathbb{Z})^2$, we iterate f and record a dictionary mapping each visited state to the step index at which it was first seen. When a revisit occurs, the cycle length is the current step index minus the recorded index. We take the maximum over all sampled initial states. For $m^2 \leq 2500$, all states are enumerated; otherwise, 500 random initial states are sampled with fixed seed.

B Trajectory Period Verification

To verify the trajectory period for seed 42, burn 300: we generate 80,000 consecutive terms starting from burn-in 0, then search for the smallest $T \in \{1, \dots, 50000\}$ such that $s_i = s_{i+T}$ for all $i \in \{0, \dots, 29\}$. Verified trajectory periods used in neural experiments:

Configuration	State-space max T	Trajectory T
$p = 5$ H-sat $m = 5$	25	25
$p = 5$ H-sat $m = 25$	212	125
$p = 5$ H-sat $m = 125$	7,295	7,295
$p = 5$ H-viol $m = 25$	30	30
$p = 7$ H-sat $m = 49$	1,083	1,083
$p = 7$ H-viol $m = 49$	46	46

Note that for $p = 5$, $m = 25$: the trajectory period (125) differs from the state-space maximum (212). The seed-42 trajectory enters a length-125 orbit rather than the maximum-length orbit. This does not materially affect the neural results because $125/6 = 20.8$ remains inside the clearly learnable regime, where the MLP reaches 100% accuracy.

C MLP Hyperparameter Sensitivity

For the H-sat $m = 125$ configuration (trajectory period 7,295, $T/L_{\text{in}} = 1216$), we tested hidden layer sizes ranging from (128) to (512, 256, 128), learning rates from 10^{-4} to 10^{-2} , and epoch counts from 50 to 500. Validation accuracy ranged from 1.1% to 4.2% across all 36 combinations (random baseline 0.8%). No combination exceeded $2\times$ random baseline. The failure is not an artifact of underfitting.

D Transformer Architecture Details

```
SeqTransformer(
  embed: Linear(1 -> 64)
  pos_enc: Embedding(6, 64)
  encoder: TransformerEncoder(
    layer0: TransformerEncoderLayer(d=64, nhead=4, ff_dim=128, dropout=0)
    layer1: TransformerEncoderLayer(d=64, nhead=4, ff_dim=128, dropout=0)
  )
  head: Linear(64*6=384 -> m)
)
Total parameters: ~250K (m=125), ~250K (m=25)
```

Optimizer: Adam, lr=1e-3, batch_size=256, epochs=50

References

- Berlekamp, E. R. (1968). *Algebraic Coding Theory*. McGraw-Hill, New York.
- Blum, M. & Micali, S. (1984). How to generate cryptographically strong sequences of pseudorandom bits. *SIAM Journal on Computing*, 13(4), 850–864.
- Gouvêa, F. Q. (1997). *p-adic Numbers: An Introduction* (2nd ed.). Springer, New York.
- Ireland, K. & Rosen, M. (1990). *A Classical Introduction to Modern Number Theory* (2nd ed.). Springer, New York.
- Knuth, D. E. (1997). *The Art of Computer Programming, Vol. 2: Seminumerical Algorithms* (3rd ed.). Addison-Wesley, Reading, MA.
- Massey, J. L. (1969). Shift-register synthesis and BCH decoding. *IEEE Transactions on Information Theory*, 15(1), 122–127.
- Harris, C. R., et al. (2020). Array programming with NumPy. *Nature*, 585, 357–362.
- Paszke, A., et al. (2019). PyTorch: An imperative style, high-performance deep learning library. *Advances in Neural Information Processing Systems*, 32.
- Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- Silverman, J. H. (2007). *The Arithmetic of Dynamical Systems*. Springer, New York.
- Vasiga, T. & Shallit, J. (2004). On the iteration of certain quadratic maps over $\text{GF}(p)$. *Discrete Mathematics*, 277(1–3), 219–240.
- Vaswani, A., et al. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30.
- Bartlett, P. L., Maass, W., & Williamson, R. C. (1998). The VC dimension of neural networks. *Handbook of Brain Theory and Neural Networks*, MIT Press.
- Hahn, M. (2020). Theoretical limitations of self-attention in neural sequence models. *Transactions of the Association for Computational Linguistics*, 8, 156–171.
- Neyshabur, B., Bhojanapalli, S., McAllester, D., & Srebro, N. (2017). Exploring generalization in deep learning. *Advances in Neural Information Processing Systems*, 30.
- Jiang, Y., et al. (2020). Fantastic generalization measures and where to find them. *International Conference on Learning Representations*.
- Olah, C., Cammarata, N., Schubert, L., Goh, G., Petrov, M., & Carter, S. (2020). Zoom in: An introduction to circuits. *Distill*, 5(3), e00024.001.
- Elhage, N., et al. (2021). A mathematical framework for transformer circuits. *Anthropic*, <https://transformer-circuits.pub/2021/framework/index.html>.
- Olsson, C., Elhage, N., Nanda, N., Joseph, N., DasSarma, N., Henighan, T., Mann, B., Askell, A., Bai, Y., Chen, A., Conerly, T., Drain, D., Ganguli, D., Hatfield-Dodds,

- Z., Hernandez, D., Johnston, S., Jones, A., Kernion, J., Lovitt, L., Ndousse, K., Amodei, D., Brown, T., Clark, J., Kaplan, J., McCandlish, S., & Olah, C. (2022). In-context learning and induction heads. *Anthropic*, <https://transformer-circuits.pub/2022/in-context-learning-and-induction-heads/index.html>.
- Orvieto, A., Smith, J., Gu, A., Fernando, A., Hyland, S., & Rubanova, Y. (2023). Resurrecting recurrent neural networks for long sequences. *International Conference on Machine Learning*.
- National Institute of Standards and Technology (2010). A Statistical Test Suite for Random and Pseudorandom Number Generators for Cryptographic Applications. *NIST Special Publication 800-22 Revision 1a*.
- Wolpert, D.H. & Macready, W.G. (1997). No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1), 67–82.
- Wei, J., et al. (2022). Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35.
- Power, A., Burda, Y., Edwards, H., Babuschkin, I., & Misra, V. (2022). Grokking: Generalization beyond overfitting on small algorithmic datasets. *arXiv preprint arXiv:2201.02177*.